

```
#if !defined(__IMAIL_H)
#define __IMAIL_H

#define IMAIL_MAX_USERS 32
#define IMAIL_NAME_LEN 32
#define IMAIL_PASS_LEN 32
#define IMAIL_SUBJ_LEN 64
#define IMAIL_BODY_LEN 384
#define IMAIL_MSG_LEN (IMAIL_NAME_LEN + IMAIL_SUBJ_LEN + IMAIL_BODY_LEN + 20)

/*
 * For simplicity, we compare the contents of this structure exactly
 * rather than treating the two fields as strings, thus the values must
 * be zero-padded before passing them into the kernel.
 */
typedef struct Imail_auth_t Imail_auth_t;
struct Imail_auth_t {
    char name[IMAIL_NAME_LEN];
    char passphrase[IMAIL_PASS_LEN];
};

/*
 * For simplicity, we compare the contents of the recipient field of this
 * structure exactly rather than treating it as a string, thus the name
 * must be zero-padded before passing it into the kernel.
 */
typedef struct Imail_header_t Imail_header_t;
struct Imail_header_t {
    char recipient[IMAIL_NAME_LEN];
    char subject[IMAIL_SUBJ_LEN];
};

/* See asm/ioctl.h for IOCTL macro definitions. */

/* authenticate a user */
#define IMAIL_AUTH_USER _IOW('X', 0x00, struct Imail_auth_t)
#define IMAIL_SET_PASS _IOW('X', 0x01, struct Imail_auth_t)
#define IMAIL_WRITE_MSG _IOW('X', 0x02, struct Imail_header_t)
#define IMAIL_DELETE_MSG _IO('X', 0x03)

/* IOCTL's performed only by Imail administrator. */
#define IMAIL_ADD_USER _IOW('X', 0x80, struct Imail_auth_t)
#define IMAIL_DELETE_USER _IOW('X', 0x81, struct Imail_auth_t)

#endif /* __IMAIL_H */
```

Iml.c

```

#include <linux/config.h>
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/types.h>
#include <linux/capability.h>
#include <linux/fs.h>
#include <linux/poll.h>
#include <linux/rwsem.h>
#include <linux/slab.h>

#include <asm/io.h>
#include <asm/uaccess.h>
#include <linux/Imail.h>

#define IMAIL_DEV_NAME "Imail" /* device name to be used in registration */

/*
 * LOCK ORDERING RULES FOR THIS MODULE (when >1 needed):
 *   user list r/w semaphore (user_list_rwsem)
 *   user data structure semaphore (sem)
 *
 * NOTE: Only one user's lock should be held at a time to avoid cycles
 *   in lock acquisition.
 */

/*
 * The Iml_msg_t is used internally to hold messages for Iml users.
 * Messages are dynamically allocated when a user requests to write one
 * (via the IMAIL_WRITE_MSG ioctl command), at which point the header
 * is filled in. Subsequent writes fill the body, and when the user
 * requests that it be sent (via fsync), closes the Iml descriptor,
 * or starts a new message, the message is delivered to the recipient.
 * If a user is deleted, the pending message is freed immediately, as
 * are any messages delivered to that user. Otherwise, message memory
 * is reclaimed when a user deletes a message (via the IMAIL_DELETE_MSG
 * ioctl command).
 */
typedef struct Iml_msg_t Iml_msg_t;
struct Iml_msg_t {
    Iml_msg_t* next; /* next message in list */
    int length; /* length of body */
    char recipient[IMAIL_NAME_LEN]; /* name of recipient */
    char body[IMAIL_MSG_LEN]; /* body of message */
};

/*
 * When a message is written/read, it is formatted as shown below:
From : Steve Lumetta
Subject: Iml message formatting
(body)

+ the name in the from line is padded with spaces to make it
  IMAIL_NAME_LEN characters long, and terminated with LF
+ the subject line is padded to IMAIL_SUBJ_LEN characters, terminated with LF
 */

/*
 * Iml_user_data_init
 * DESCRIPTION: Initialize the data for a new user; does NOT initialize
 * authentication data.
 */

```

```

 * Data structure for each Iml user. A single structure is statically
 * allocated and cannot be deleted--this structure holds the data for
 * the Iml administrator. All other user structures are dynamically
 * allocated and form a linked list. Access to the list is protected
 * by a reader/writer semaphore, user_list_rwsem.
 */
typedef struct Iml_user_data_t Iml_user_data_t;
struct Iml_user_data_t {
    /*
     * This part of the data structure (including the list structure)
     * is synchronized with user_list_rwsem.
     */
    Iml_user_data_t* next; /* link to next user */
    Iml_auth_t user_auth; /* user authentication data */

    /* Wait queues contain their own synchronization variables. */
    wait_queue_head_t read_queue; /* wait queue for new messages */

    /*
     * The semaphore below is used to synchronize access to the rest
     * of the data structure. Reading the values of active_file and
     * msg_list is allowed without the semaphore (to simplify the
     * implementation of waiting for messages), but dereferencing them
     * is not. Note that checking for user deletion (active_file is NULL
     * while the files private_data points to this structure) MUST be
     * done with the semaphore to avoid races with the deletion code.
     */
    struct semaphore sem; /* for user data synchronization */
    struct file* active_file; /* non-NULL when authenticated */
    Iml_msg_t* msg_list; /* mailbox (linked list) */
    Iml_msg_t* writing; /* message being written */

};

/* module entry/exit and system call declarations */
static int Iml_init(void);
static void Iml_exit(void);
static ssize_t Iml_read(struct file*, char*, size_t, loff_t*);
static ssize_t Iml_write(struct file*, const char*, size_t, loff_t*);
static unsigned int Iml_poll(struct file*, struct poll_table_struct*);
static int Iml_ioctl(struct inode*, struct file*, unsigned int,
                    unsigned long);
static int Iml_release(struct inode*, struct file*);
static int Iml_fsync(struct file*, struct dentry*, int datasync);

/* module system call jump table */
static struct file_operations Iml_fops;

/* static variables for module */
static kdev_t Iml_dev; /* assigned major device number */
static Iml_user_data_t user_list; /* user list head (administrator) */
static DECLARE_RWSEM (user_list_rwsem); /* user list r/w semaphore */

/*-----*
 * The following functions are basic operations and tests on Iml user
 * data structures. They appear first to allow inlining by the compiler.
 *-----*/

/*
 * Iml_user_data_init
 * DESCRIPTION: Initialize the data for a new user; does NOT initialize
 * authentication data.
 */

```

Iemail.c

```

* INPUTS: udata -- an uninitialized user
* OUTPUTS: udata -- the newly initialized user
* RETURN VALUE: none
* SIDE EFFECTS: none
* SYNCHRONIZATION: User must NOT be accessible from user list during
  this function; link user into list AFTER
  initialization.
*/
static void
Iemail_user_data_init (Iemail_user_data_t* udata)
{
    udata->next = NULL;
    init_MUTEX (&udata->sem);
    udata->active_file = NULL;
    udata->msg_list = NULL;
    udata->writing = NULL;
    init_waitqueue_head (&udata->read_queue);
}

/* Iemail_user_data_free
 * DESCRIPTION: Free all dynamically allocated data associated with a
 *              user (but not the user structure).
 * INPUTS: udata -- the user structure
 * OUTPUTS: none
 * RETURN VALUE: none
 * SIDE EFFECTS: frees all messages from list and writing
 * SYNCHRONIZATION: User and/or messages must NOT be accessible to other
 *                  code during this function.
 */
static void
Iemail_user_data_free (Iemail_user_data_t* udata)
{
    Iemail_msg_t* msg;    /* message to be freed */

    /* Free all messages in message list. */
    while (NULL != (msg = udata->msg_list)) {
        udata->msg_list = msg->next;
        kfree (msg);
    }

    /* Discard any pending message and issue a warning--message should
     * have been sent or deleted before this point.
     */
    if (NULL != udata->writing) {
        printk ("Iemail deleting unsent message.\n");
        kfree (udata->writing);
    }
}

/* Iemail_user_is_administrator
 * DESCRIPTION: Checks whether a specified user is the administrator.
 *              Any caller with system administrator privileges also
 *              has Iemail administrative privileges. Otherwise, the
 *              caller must have authenticated as the Iemail administrator
 *              (and udata points to the static element at the head of
 *              the user list).
 * INPUTS: udata -- the user structure, or NULL if not authenticated
 * OUTPUTS: none
 * RETURN VALUE: 1 if the user has administrative privileges, 0 otherwise
 * SIDE EFFECTS: none

```

```

* SYNCHRONIZATION: none necessary; udata pointer is not dereferenced
*/
static int
Iemail_user_is_administrator (const Iemail_user_data_t* udata)
{
    /*
     * Either the Iemail administrator or the system administrator.
     * Note that we check system administrator capability REGARDLESS
     * of the state of the calling process, so a system administrator
     * can be unauthenticated or even authenticated as a user and
     * subsequently deleted and still exercise administrator privileges
     * with the same open file descriptor.
     */
    return (udata == &user_list || capable (CAP_SYS_ADMIN));
}

/* Iemail_user_auth_matches
 * DESCRIPTION: Checks whether an authorization structure provided
 *              matches that of a given user.
 * INPUTS: udata -- the user structure to be checked against
 *          auth -- the authentication data provided
 * OUTPUTS: none
 * RETURN VALUE: 1 if the data match, 0 otherwise
 * SIDE EFFECTS: none
 * SYNCHRONIZATION: Acquire a lock (read or write) on the user list
 *                  before calling this function.
 */
static int
Iemail_user_auth_matches (const Iemail_user_data_t* udata,
                          const Iemail_auth_t* auth)
{
    return (0 == memcmp (&udata->user_auth, auth, sizeof (*auth)));
}

/* Iemail_user_name_matches
 * DESCRIPTION: Checks whether a name provided matches that of a
 *              given user.
 * INPUTS: udata -- the user structure to be checked against
 *          name -- the name provided
 * OUTPUTS: none
 * RETURN VALUE: 1 if the name matches, 0 otherwise
 * SIDE EFFECTS: none
 * SYNCHRONIZATION: Acquire a lock (read or write) on the user list
 *                  before calling this function.
 */
static int
Iemail_user_name_matches (const Iemail_user_data_t* udata,
                          const char name[IEMAIL_NAME_LEN])
{
    return (0 == memcmp (&udata->user_auth.name, name, IEMAIL_NAME_LEN));
}

/*-----
 * The following two functions are used by the kernel to install and
 * uninstall the Iemail module.
 *-----
 */
/* Iemail_init
 * DESCRIPTION: Initialize the Iemail module.  Creates a single

```

Iemail.c

```

* user, who is also the administrator.
*
* INPUTS: none
* OUTPUTS: none
* RETURN VALUE: 0 on success, error number from register_chrdev on failure
* SIDE EFFECTS: registers the Iemail pseudo-device as a character device
*               with dynamically-allocated major number; prints a message
*               with the allocated major number
* * SYNCHRONIZATION: called atomically by kernel, using the big kernel lock
*               to prevent devices from being opened before this code
*               finishes
*/
static int __init
Iemail_init ()
{
    int rval;
    int i;

    Iemail_user_data_init (&user_list);
    memset (&user_list.user_auth, 0, sizeof (user_list.user_auth));
    strcpy (user_list.user_auth.name, "administrator");
    strcpy (user_list.user_auth.passphrase, "ece398");
    if ((rval = register_chrdev (0, IMAIL_DEV_NAME, &Iemail_fops)) < 0)
        return rval;

    Iemail_dev = MKDEV (rval, 0);
    printk ("Registered Iemail as device %d.\n", rval);

    return 0;
}

/*
* Iemail_exit
* DESCRIPTION: Remove the Iemail module. Frees all dynamically allocated
*               memory associated with the module.
*
* INPUTS: none
* OUTPUTS: none
* RETURN VALUE: none
* SIDE EFFECTS: destroys user list and messages
* SYNCHRONIZATION: The kernel is responsible for ensuring that no files
*               are actively using the module (linux-2.4 and later).
*/
static void
Iemail_exit ()
{
    Iemail_user_data_t* udata; /* user data to be freed */

    /* Free any messages in administrator's data. */
    Iemail_user_data_free (&user_list);

    /* Walk through list of users, freeing each one's data. */
    while (user_list.next != NULL) {
        udata = user_list.next;
        Iemail_user_data_free (udata);
        user_list.next = udata->next;
        kfree (udata);
    }

    /*
    * If the unregister call fails, we can't do anything except panic
    * and bring down the kernel, which seems like overkill.
    */
    (void)unregister_chrdev (MAJOR (Iemail_dev), IMAIL_DEV_NAME);
}

```

```

/*-----*
* The functions below are helper functions for the Iemail system calls. *
*-----*/

/*
* Iemail_find_user
* DESCRIPTION: Find a user by name.
* INPUTS: name -- the name of the user sought
* OUTPUTS: none
* RETURN VALUE: the user's data, or NULL if no such user exists
* SIDE EFFECTS: none
* SYNCHRONIZATION: Acquire a lock (read or write) on the user list
*               before calling this function.
*/
static Iemail_user_data_t*
Iemail_find_user (const char name[IMAIL_NAME_LEN])
{
    Iemail_user_data_t* udata; /* user data structure for search */

    /* Attempt to find the requested user in the list. */
    for (udata = &user_list; NULL != udata; udata = udata->next) {

        /* Compare the name provided with that of the user. */
        if (Iemail_user_name_matches (udata, name))
            return udata;
    }

    /* No luck: return NULL. */
    return NULL;
}

/*
* Iemail_grab_pending
* DESCRIPTION: Extract a pending message out of a user's data structure
*               for delivery or discard (in the case of user deletion).
*
* INPUTS: udata -- the user's data
* OUTPUTS: udata -- the user's data after message is removed
* RETURN VALUE: the message, or NULL if none was being written
* SIDE EFFECTS: none
* SYNCHRONIZATION: Acquire the user data semaphore (udata->sem)
*               before calling this function.
*/
static Iemail_msg_t*
Iemail_grab_pending (Iemail_user_data_t* udata)
{
    Iemail_msg_t* pending; /* the pending message */

    pending = udata->writing;
    udata->writing = NULL;
    return pending;
}

/*
* Iemail_deliver
* DESCRIPTION: Deliver a message to its intended recipient, appending
*               it to the end of the recipient's message list.
* INPUTS: msg -- the message to be delivered
*               have_list_lock -- non-zero if the caller has a lock (read or
*               write) on the user list
* OUTPUTS: none
* RETURN VALUE: 0 on success, or -ENOENT if recipient does not exist

```

```

* SIDE EFFECTS: message appended to recipient's message list
* SYNCHRONIZATION: May be called with or without a user list lock.
*
* The second argument specifies the state on entry.
* If no lock is held already, a read lock is acquired
* on the user list. If the recipient exists, the
* user data semaphore for the recipient is acquired for
* appending the message to the recipient's message list.
*/
static int
Iemail_deliver (Iemail_msg_t* msg, int have_user_list_lock)
{
    Iemail_user_data_t* udata; /* recipient's user data structure */
    Iemail_msg_t** find; /* used to find end of message list */

    /* Lock the user list to find the recipient. */
    if (!have_user_list_lock)
        down_read (&user_list_rwlock);

    /* Find the user based on the name in the message. If no recipient
     * is found, throw it away and return an error.
     */
    if (NULL == (udata = Iemail_find_user (msg->recipient))) {
        up_read (&user_list_rwlock);
        kfree (msg);
        return -ENOENT;
    }

    /* Lock the user's data, add the message to the end of the user's
     * list, wake up any programs waiting for new messages for this
     * user, and unlock the data.
     */
    down (&udata->sem);
    for (find = &udata->msg_list; NULL != (*find); find = &(*find)->next);
    *find = msg;
    wake_up_interruptible (&udata->read_queue);
    up (&udata->sem);

    /* Release the user list lock if acquired by this function. */
    if (!have_user_list_lock)
        up_read (&user_list_rwlock);

    /* Return success. */
    return 0;
}

/* nul_to_space_cpy
 * DESCRIPTION: Copy bytes from one place to another, transforming ASCII
 * NUL characters into spaces. Used in formatting messages.
 * INPUTS: src -- pointer to source buffer
 * len -- number of bytes to copy
 * OUTPUTS: dst -- destination buffer
 * RETURN VALUE: none
 * SIDE EFFECTS: none
 * SYNCHRONIZATION: none required per se
 */
static void
nul_to_space_cpy (char* dst, const char* src, size_t len)
{
    char c; /* character to copy */

    /* Copy characters one by one. */

```

Iemail.c

```

while (len-- > 0) {
    /* Transform NUL to SPACE */
    if ('\0' == (c = *src++))
        c = ' ';
    *dst++ = c;
}

/*-----
 * The functions below are the Iemail system calls. The ioctl commands are *
 * each implemented as a separate function.
 *-----*/

/* Iemail_read
 * DESCRIPTION: Implements the read system call for Iemail. Data are
 * read from the first message in the user's list (only).
 * Copies data to buffer and advances file pointer.
 * INPUTS: f -- the file structure used for reading
 * buf -- the user's buffer
 * len -- the length of the buffer
 * offp -- pointer to the file offset
 * OUTPUTS: buf -- buffer with data copied from first message
 * offp -- updated file pointer
 * RETURN VALUE: number of bytes read on success
 * -EPERM if user not authenticated or deleted after
 * authentication
 * -EAGAIN when no message available (non-blocking I/O only)
 * -EINTR when signal received
 * -EFAULT if buffer is bad
 * SIDE EFFECTS: advances file pointer; may sleep
 * SYNCHRONIZATION: Obtains the associated user's user data semaphore.
 * May sleep on user data's read_queue waiting for a
 * message to arrive.
 */
static ssize_t
Iemail_read (struct file* f, char* buf, size_t len, loff_t* offp)
{
    Iemail_user_data_t* udata; /* user data structure */
    Iemail_msg_t* msg; /* message to be read */
    ssize_t n_read; /* bytes read */
    loff_t off; /* offset */

    /* If user has not authenticated yet, permission is denied. */
    if (NULL == (udata = f->private_data))
        return -EPERM;

    /* Loop to wait for data to be available for reading. */
    while (1) {
        /* Start by obtaining user data lock and checking for user deletion. */
        down (&udata->sem);
        if (NULL == udata->active_file) {
            up (&udata->sem);
            return -EPERM;
        }

        /* Have data to read? Break out of this loop. */
        if (NULL != (msg = udata->msg_list))
            break;

        /* No data yet. Release semaphore and either return failure (for

```

Imail.c

```

/*
 * non-blocking I/O) or wait for a message or user deletion.
 */
up (&udata->sem);
if (0 != (f->f_flags & O_NONBLOCK))
    return -EAGAIN;
if (wait_event_interruptible
    (udata->read_queue, (NULL == udata->active_file ||
    NULL != udata->msg_list)))
    return -PRESTARTSYS; /* got a signal */
/*
 * ERESTARTSYS prevents use of signals to break out of system calls.
 * The semantics are taken from BSD. To change this property, use
 * sigaction and turn off SA_RESTART for the signal used to kick
 * system calls out of blocking.
 */
}

/*
 * We have an undeleted user, a message to read, and a semaphore to
 * protect us. Note that the default kernel seek functions do not
 * allow file offsets < 0.
 */
n_read = 0;
off = *offp;
if (msg->length > off) {
    /* There are bytes left to read--read as many as fit in buffer. */
    if (len < (n_read = msg->length - off))
        n_read = len;

    /* Copy the bytes into the user's buffer. */
    if (copy_to_user (buf, msg->body + off, n_read))
        n_read = -EFAULT;
    else
        (*offp) = off + n_read;
}

/*
 * Release the semaphore and return the number of bytes read
 * (or error number).
 */
up (&udata->sem);
return n_read;
}

/*
 * Imail_write
 * DESCRIPTION: Implements the write system call for Imail. Data are
 * written to a message allocated with ioctl. Copies data
 * to message and increases message length.
 * INPUTS: f -- the file structure used for reading
 *         buf -- the user's buffer
 *         len -- the length of the buffer
 *         offp -- IGNORED
 * OUTPUTS: none
 * RETURN VALUE: number of bytes written on success
 *              -EPERM if user not authenticated or deleted after
 *                  authentication
 *              -EIO when no message is being written currently
 *              -EBIG when no space is left in the current message
 *              -EFAULT if buffer is bad
 * SIDE EFFECTS: updates message body length
 * SYNCHRONIZATION: Obtains the associated user's user data semaphore.
 *                  May sleep with semaphore if user buffer is not
 *                  in memory (page fault).
 */
static ssize_t
Imail_write (struct file* f, const char* buf, size_t len, loff_t* offp)
{
    Imail_user_data_t* udata; /* user data structure */
    Imail_msg_t* msg; /* message to be written */
    ssize_t n_written; /* bytes written */

    /* If user has not authenticated yet, permission is denied. */
    if (NULL == (udata = f->private_data))
        return -EPERM;

    /* Get the user data semaphore. */
    down (&udata->sem);

    /* Make checks, then write. */
    n_written = 0;
    if (NULL == udata->active_file) {
        /* User was deleted after authentication. */
        n_written = -EPERM;
    } else if (NULL == (msg = udata->writing)) {
        /* User has not started writing a message. */
        n_written = -EIO;
    } else if (0 >= (n_written = (IMAIL_MSG_LEN - msg->length))) {
        /* No space left to write. */
        n_written = -EBIG;
    } else {
        /* Limit to characters provided. */
        if (len < n_written)
            n_written = len;

        /* Copy the bytes from the user's buffer. */
        if (copy_from_user (msg->body + msg->length, buf, n_written))
            n_written = -EFAULT;
        else
            msg->length += n_written;
    }

    /* Release the semaphore and return the number of bytes written
     * (or error number).
     */
    up (&udata->sem);
    return n_written;
}

/*
 * Imail_poll
 * DESCRIPTION: Implements the poll and select system calls for Imail.
 * Imail pseudo-devices are always writable if write
 * permission is held, and readable if read permission is
 * held and a message is available.
 * INPUTS: f -- the file structure used for reading
 *         ptab -- the poll table used to collect wait queues
 * OUTPUTS: none
 * RETURN VALUE: mask of permitted operations on success
 *              POLLERR if user not authenticated or deleted after
 *                  authentication
 * SIDE EFFECTS: adds user's read queue to the poll table
 * SYNCHRONIZATION: The user data semaphore is acquired only to check for
 *                  user deletion, although user's msg_list is read
 *                  (not dereferenced) for read permission check.
 *                  Associating our read queue with the poll table
 *                  before the check should suffice to handle race
 */

```

```

*
* conditions: if message is delivered after poll_wait,
* poll should not sleep; messages delivered before
* poll wait (and not deleted before check) are seen
* by check.
*/
static unsigned int
Iml_poll (struct file* f, struct poll_table_struct* ptab)
{
    Iml_user_data_t* udata; /* user data structure */
    unsigned int mask; /* return value */

    /* If user has not authenticated yet, return an error. */
    if (NULL == (udata = f->private_data))
        return POLLERR;

    /* Tell poll to wake up if our user's read queue is signalled. */
    poll_wait (f, &udata->read_queue, ptab);

    /*
     * Need the semaphore to check whether user was deleted.
     * If so, poll will remove our wait queue from the table
     * before returning, and udata will not be deleted before
     * poll returns.
     */
    down (&udata->sem);
    if (NULL == udata->active_file) {
        up (&udata->sem);
        return POLLERR;
    }
    up (&udata->sem);

    /* Calculate mask of available operations. */
    mask = 0;
    if (0 != (f->f_mode & FMODE_READ) && NULL != udata->msg_list)
        mask |= (POLLIN | POLLRDNORM);
    if (0 != (f->f_mode & FMODE_WRITE))
        mask |= (POLLOUT | POLLWRNORM);

    return mask;
}

/*
 * Iml_authenticate
 * DESCRIPTION: Authenticate a user based on the name and password provided.
 * Only unauthenticated users are allowed to authenticate.
 * Those that have been booted off via user deletion must
 * close all descriptors to ensure atomic deletion of the user
 * data structure. Reauthentication is technically possible,
 * but introduces some complexity in the sense that any
 * operations on the old user must be switched over to the
 * new user or cancelled. It's cleaner to disallow
 * reauthentication.
 *
 * INPUTS: f -- file structure on which ioctl was called
 * arg -- a pointer to the user's authentication data
 * OUTPUTS: none
 * RETURN VALUE: 0 on success
 *              -EPERM if file was already authenticated, deleted after
 *                  authentication
 *              -EFAULT if the authentication data buffer was bad
 *              -EBUSY if the requested user is in use by another file
 *              -EACCES if authentication data match no user in list
 * SIDE EFFECTS: points f->private_data to the user data structure if
 * authentication succeeds
 * SYNCHRONIZATION: Obtains a read lock on the user list. If authentication

```

```

*
* succeeds for some user, obtains that user's user data
* semaphore.
*/
static int
Iml_authenticate (struct file* f, void* arg)
{
    Iml_auth_t auth; /* authentication data */
    Iml_user_data_t* udata; /* user data structure */
    int rval; /* return value */

    /* Authenticated already or been booted? Close and re-open. */
    if (NULL != (udata = f->private_data))
        return -EPERM;

    /* Read authentication structure into kernel address space. */
    if (copy_from_user (&auth, arg, sizeof (auth)))
        return -EFAULT;

    /* We need to read authentication data from the user list. */
    down_read (&user_list_rwlock);

    /* Attempt to find the requested user in the list. */
    for (udata = &user_list; NULL != udata; udata = udata->next) {
        /* Compare the authentication data provided with those of
         * the user.
         */
        if (Iml_user_auth_matches (udata, &auth)) {

            /* We have a match! Try to install this file as the active one. */
            down (&udata->sem);

            /* Check if already in use. */
            if (NULL != udata->active_file)
                rval = -EBUSY;
            else {
                /* Record the user information in both directions. */
                udata->active_file = f;
                f->private_data = udata;
                rval = 0;
            }

            /* We're done. */
            up (&udata->sem);
            up_read (&user_list_rwlock);
            return rval;
        }
    }

    /* No luck in authenticating; permission denied! */
    up_read (&user_list_rwlock);
    return -EACCES;
}

/* Iml_set_pass
 * DESCRIPTION: Set a user's password. Changing a user's password
 * requires that the entity making the change either
 * authenticate first as that user or have Iml administrative
 * privileges.
 * INPUTS: f -- file structure on which ioctl was called
 * arg -- a pointer to a user's name and new password
 * OUTPUTS: none
 * RETURN VALUE: 0 on success

```

```

* -EFAULT if the authentication data buffer was bad
* -ENOENT if the named user does not exist
* -EPERM if the attempt to change was illegal
*
* SIDE EFFECTS: may change authentication data for a user
* SYNCHRONIZATION: Obtains a write lock on the user list.
* SECURITY: Note that the error values in this function confirm user
*          existence. If this type of leak is undesirable, return
*          -EPERM instead of -ENOENT.
*/
static int
Imail_set_pass (struct file* f, void* arg)
{
    Imail_auth_t auth; /* authentication data */
    Imail_user_data_t* udata; /* user data structure */
    Imail_user_data_t* target; /* target user data structure */
    int rval; /* return value */

    /* Read authentication structure into kernel address space. */
    if (copy_from_user (&auth, arg, sizeof (auth)))
        return -EFAULT;

    /*
     * We may need to change authentication data, so we write lock the
     * user list.
     */
    down_write (&user_list_rwlock);

    /* Assume success. */
    rval = 0;
    udata = f->private_data;

    /* Find the user by name; if no match is found, return an error. */
    if (NULL == (target = Imail_find_user (auth.name))) {
        rval = -ENOENT;
    } else if (!Imail_user_is_administrator (udata) && udata != target) {
        /*
         * Changing a password for any user but yourself requires
         * administrator privileges. Keep in mind that the file f
         * may be unauthenticated, in which case f->private_data is
         * NULL, but the caller is the system administrator.
         */
        rval = -EPERM;
    } else {
        /* Change the password by overwriting the authentication structure. */
        memcpy (&target->user_auth, &auth, sizeof (auth));
    }

    /* All done. Release the lock and return. */
    up_write (&user_list_rwlock);
    return rval;
}

/*
 * DESCRIPTION: Prepare to write an Imail message. Allocate a message
 * buffer and attach it to the user's data structure. If
 * the user was already writing a message, send it to its
 * recipient.
 *
 * INPUTS: f -- file structure on which ioctl was called
 *          arg -- a pointer to a message header
 *
 * OUTPUTS: none
 *
 * RETURN VALUE: 0 on success
 *               -EPERM if user not authenticated or deleted after
 *                 authentication
 */
Imail_write_msg (struct file* f, void* arg)
{
    static int
    Imail_write_msg (struct file* f, void* arg)
    {
        Imail_header_t header; /* header for new message */
        Imail_user_data_t* udata; /* user data structure */
        Imail_msg_t* msg; /* message to be delivered */
        int rval; /* return value */
        char* wpos; /* write position in message body */

        /* Initialize variables an check for user authentication. */
        if (NULL == (udata = f->private_data))
            return -EPERM;
        msg = NULL;
        rval = 0;

        /* Check that the file is writable. */
        if (0 == (f->f_mode & FMODE_WRITE))
            return -EBADF;

        /* Read header structure into kernel address space. */
        if (copy_from_user (&header, arg, sizeof (header)))
            return -EFAULT;

        /* Check for recipient. Recipient MAY be deleted between this check
         * and message creation, but could also be deleted at any point between
         * this call and the call that delivers the message. This check works
         * for the more common case of writing to someone who has never existed.
         */
        down_read (&user_list_rwlock);
        if (NULL == Imail_find_user (header.recipient))
            rval = -ENOENT;
        up_read (&user_list_rwlock);

        /*
         * Lock the user data, check whether this user was deleted after
         * authentication, extract any pending messages for delivery, and
         * allocate a new message for writing. In any case, unlock at the end.
         */
        down (&udata->ssem);
        /* Check for user deletion. */
        if (NULL == udata->active_file)
            rval = -EPERM; /* overrides absent recipient error */
        else if (rval == 0) {
            /* does not create message unless recipient existed */
            msg = Imail_grab_pending (udata);
            udata->writing = kmalloc (sizeof (Imail_msg_t), GFP_KERNEL);
            if (NULL == udata->writing)
                rval = -ENOMEM;
            else {
                udata->writing->next = NULL;
                memcpy (&udata->writing->recipient, header.recipient,
                        sizeof (header.recipient));
            }
        }
    }
}

```

```

* -EBADF if write permission not held on file
* -EFAULT if the authentication data buffer was bad
* -ENOMEM if the kernel runs out of memory (!)
* -ENOENT if a previous message could not be delivered
*
* SIDE EFFECTS: delivers any message currently being written
* SYNCHRONIZATION: Obtains a user list read lock to check for existence
*                  of intended recipient. Obtains the associated user's
*                  user data semaphore. If message must be delivered,
*                  obtains both a user list read lock and the recipient's
*                  semaphore.
*/
static int
Imail_write_msg (struct file* f, void* arg)
{
    Imail_header_t header; /* header for new message */
    Imail_user_data_t* udata; /* user data structure */
    Imail_msg_t* msg; /* message to be delivered */
    int rval; /* return value */
    char* wpos; /* write position in message body */

    /* Initialize variables an check for user authentication. */
    if (NULL == (udata = f->private_data))
        return -EPERM;
    msg = NULL;
    rval = 0;

    /* Check that the file is writable. */
    if (0 == (f->f_mode & FMODE_WRITE))
        return -EBADF;

    /* Read header structure into kernel address space. */
    if (copy_from_user (&header, arg, sizeof (header)))
        return -EFAULT;

    /* Check for recipient. Recipient MAY be deleted between this check
     * and message creation, but could also be deleted at any point between
     * this call and the call that delivers the message. This check works
     * for the more common case of writing to someone who has never existed.
     */
    down_read (&user_list_rwlock);
    if (NULL == Imail_find_user (header.recipient))
        rval = -ENOENT;
    up_read (&user_list_rwlock);

    /*
     * Lock the user data, check whether this user was deleted after
     * authentication, extract any pending messages for delivery, and
     * allocate a new message for writing. In any case, unlock at the end.
     */
    down (&udata->ssem);
    /* Check for user deletion. */
    if (NULL == udata->active_file)
        rval = -EPERM; /* overrides absent recipient error */
    else if (rval == 0) {
        /* does not create message unless recipient existed */
        msg = Imail_grab_pending (udata);
        udata->writing = kmalloc (sizeof (Imail_msg_t), GFP_KERNEL);
        if (NULL == udata->writing)
            rval = -ENOMEM;
        else {
            udata->writing->next = NULL;
            memcpy (&udata->writing->recipient, header.recipient,
                    sizeof (header.recipient));
        }
    }
}

```



```

/* Write a header into the message body. */
udata->writing-length = 20 + sizeof (udata->user_auth.name) +
    sizeof (header.subject);
wpos = udata->writing->body;
memcpy (wpos, "From : ", 9);
wpos += 9;
nul_to_space_cpy (wpos, udata->user_auth.name,
    sizeof (udata->user_auth.name));
wpos += sizeof (udata->user_auth.name);
memcpy (wpos, "\nSubject: ", 10);
wpos += 10;
nul_to_space_cpy (wpos, header.subject, sizeof (header.subject));
wpos += sizeof (header.subject);
memcpy (wpos, "\n", 1);
    }
}
up (&udata->sem);

/* Deliver any pending message from the user. If we are already
 * returning an error, we ignore errors in delivery. Otherwise,
 * we report them.
 */
if (NULL != msg) {
    if (rval == 0)
        rval = Imail_deliver (msg, 0);
    else
        (void)Imail_deliver (msg, 0);
}

return rval;
}

/* Imail_delete_msg
 * DESCRIPTION: Delete the first message in a user's list of messages.
 *              Deleting incoming messages is part of reading, so this
 *              command requires read permission on the file.
 * INPUTS: f -- file structure on which ioctl was called
 *          arg -- a pointer to a message header
 * OUTPUTS: none
 * RETURN VALUE: 0 on success
 *              -EPERM if user not authenticated or deleted after
 *                  authentication
 *              -EADF if read permission not held on file
 * SIDE EFFECTS: frees the memory associated with the deleted message
 * SYNCHRONIZATION: Obtains the associated user's user data semaphore.
 */
static int
Imail_delete_msg (struct file* f, void* arg)
{
    Imail_user_data_t* udata; /* user data structure */
    Imail_msg_t* discard; /* message to be deleted */
    int rval; /* return value */

    /* Initialize variables an check for user authentication. */
    if (NULL == (udata = f->private_data))
        return -EPERM;
    discard = NULL;
    rval = 0;

    /* Check that the file is readable. */
    if (0 == (f->f_mode & FMODE_READ))
        return -EADF;
}

```

```

/* Lock the user data and remove the first message from
 * the message list. If the user was deleted, we return an error.
 */
down (&udata->sem);
if (NULL == udata->active_file)
    rval = -EPERM;
else {
    if (NULL != (discard = udata->msg_list))
        udata->msg_list = discard->next;
    f->f_pos = 0; /* used to track read position in message */
}
up (&udata->sem);

/* Free the memory outside of the critical section. */
if (NULL != discard)
    kfree (discard);

return rval;
}

/* Imail_add_user
 * DESCRIPTION: Add a new user to the list. Adding a user requires
 *              that the caller have Imail administrative privileges.
 * INPUTS: f -- file structure on which ioctl was called
 *          arg -- a pointer to the new user's name and password
 * OUTPUTS: none
 * RETURN VALUE: 0 on success
 *              -EPERM if no administrative privileges are held
 *              -EFAULT if the authentication data buffer was bad
 *              -EXIST if the named user already exists
 *              -ENOMEM if the kernel runs out of memory (!)
 * SIDE EFFECTS: may allocate new user structure and insert into user list
 * SYNCHRONIZATION: Obtains a write lock on the user list.
 */
static int
Imail_add_user (struct file* f, void* arg)
{
    Imail_user_data_t* udata; /* user data structure */
    int rval; /* return value */
    Imail_auth_t auth; /* new user's authentication data */

    /* Assume success. */
    rval = 0;

    /* This command is only allowed to administrators. */
    if (!Imail_user_is_administrator (f->private_data))
        return -EPERM;

    /* Read authentication structure into kernel address space. */
    if (copy_from_user (&auth, arg, sizeof (auth)))
        return -EFAULT;

    /* Write lock the user list. */
    down_write (&user_list_rwlock);

    /* Try to add the user. */
    if (NULL != Imail_find_user (auth.name)) {
        /* User exists; we can't add another with the same name. */
        rval = -EXIST;
    } else {
        /* Allocate a user data block. */

```

Iemail.c

```

    udata = kcalloc (sizeof (Iemail_user_data_t), GFP_KERNEL);
    if (NULL == udata)
        rval = -ENOMEM;
    else {
        /* Copy the authentication data into place. */
        memcpy (&udata->user_auth, &auth, sizeof (auth));

        /* Initialize the rest of the structure. */
        Iemail_user_data_init (udata);

        /* And link the user into place after the administrator. */
        udata->next = user_list.next;
        user_list.next = udata;
    }

    /* Release our lock and return the right value. */
    up_write (&user_list_rwlock);
    return rval;
}

/* Iemail_delete_user
DESCRIPTION: Delete a user from the list. Deleting a user requires
that the caller have Iemail administrative privileges.
INPUTS: f -- file structure on which ioctl was called
arg -- a pointer to the new user's name and password
OUTPUTS: none
RETURN VALUE: 0 on success
-EPERM if no administrative privileges are held, or if
attempt is made to delete the Iemail administrator
-DEFAULT if the authentication data buffer was bad
-ENOENT if the named user does not exist
SIDE EFFECTS: may unlink user data structure from list; may free
user data structure memory
SYNCHRONIZATION: Obtains a write lock on the user list. If user
exists and is active, obtains user's user data
semaphore and marks as booted from system rather
than freeing structure. Structure is then freed
by Iemail_release.
*/
static int
Iemail_delete_user (struct file* f, void* arg)
{
    Iemail_user_data_t* udata; /* user data structure
int rval; /* return value
Iemail_auth_t auth; /* new user's authentication data
Iemail_user_data_t** find; /* loop index for user removal

    /* Assume success. */
    rval = 0;

    /* This command is only allowed to administrators. */
    if (!Iemail_user_is_administrator (f->private_data))
        return -EPERM;

    /* Read authentication structure into kernel address space. */
    if (copy_from_user (&auth, arg, sizeof (auth)))
        return -EFAULT;

    /* Write lock the user list. */
    down_write (&user_list_rwlock);

    /* Try to find the user. */
    if (NULL == (udata = Iemail_find_user (auth.name))) {
        /* User does not exist. */
        rval = -ENOENT;
    } else if (udata == &user_list) {
        /* Deletion of administrator is not permitted. */
        rval = -EPERM;
    } else {
        /* Discard user's messages and remove them from the user list. */
        Iemail_user_data_free (udata);
        for (find = &user_list.next; (*find) != udata; find = &(*find)->next);
        *find = udata->next;

        if (NULL == udata->active_file) {
            /* Not active--just delete! */
            kfree (udata);
        } else {
            /* User is active. Boot them off indirectly. */
            udata->active_file = NULL;
            wake_up_interruptible (&udata->read_queue);
        }
    }

    /* Release our lock and return the right value. */
    up_write (&user_list_rwlock);
    return rval;
}

/* Iemail_ioctl
DESCRIPTION: Implements the ioctl system call for Iemail. Specific
effects depend on the command requested (each has its
own function).
INPUTS: idx -- the index node associated with the file
f -- the file structure
cmd -- the ioctl command
arg -- operand for the command
OUTPUTS: see command functions
RETURN VALUE: -ENOENT for unknown command
other values returned directly from command function
SIDE EFFECTS: see command functions
SYNCHRONIZATION: see command functions
*/
static int
Iemail_ioctl (struct inode* idx, struct file* f, unsigned int cmd,
              unsigned long arg)
{
    /*
    * Authentication is required for all most commands, while some
    * require Iemail administrator privileges. System administrator
    * privileges suffice for the administrative commands (with or
    * without authentication as the Iemail administrator). User-level
    * commands require authentication, since they operate on the
    * user data structure. See command functions for more detail.
    */

    /* Switch on command. */
    switch (cmd) {
        case IMAIL_AUTH_USER: return Iemail_authenticate (f, (void*)arg);
        case IMAIL_SET_PASS: return Iemail_set_pass (f, (void*)arg);
        case IMAIL_WRITE_MSG: return Iemail_write_msg (f, (void*)arg);
        case IMAIL_DELETE_MSG: return Iemail_delete_msg (f, (void*)arg);
        case IMAIL_ADD_USER: return Iemail_add_user (f, (void*)arg);
        case IMAIL_DELETE_USER: return Iemail_delete_user (f, (void*)arg);
        default:
            return -ENOENT; /* invalid command */
    }
}

```

Iemail.c

```

    }
}

/*
 * Iemail_release
 * DESCRIPTION: Closes an Iemail file for the last time. Release is only
 *              called when the last operation finishes on a file, so we
 *              know that nothing else is still working with this
 *              particular file.
 * INPUTS: idx -- the index node associated with the file
 *          f -- the file structure
 * OUTPUTS: none
 * RETURN VALUE: 0 always; ignored by kernel
 * SIDE EFFECTS: delivers an pending message to its recipient; frees
 *               user data structure if user was deleted after
 *               authentication
 * SYNCHRONIZATION: Obtains the associated user's user data semaphore
 *                  in order to extract the message. Obtains both a
 *                  user list read lock and the recipient's semaphore
 *                  in order to deliver a message.
 */
static int
Iemail_release (struct inode* idx, struct file* f)
{
    Iemail_user_data_t* udata; /* user data structure */
    int i; /* loop index for finding user */
    Iemail_msg_t* msg; /* message pending from user */

    /*
     * Initialize variables. If file's owner had not authenticated yet,
     * we do nothing.
     */
    if (NULL == (udata = f->private_data))
        return 0;
    msg = NULL;

    /*
     * We need a lock on the user data to prevent another file from
     * slipping in and authenticating for this user while we're
     * cleaning up.
     */
    down (&udata->sem);

    /* Were we properly authenticated, or booted via user deletion? */
    if (NULL == udata->active_file) {
        /* Booted: delete user data. The lock is being deleted, too. */
        kfree (udata);
    } else {
        /*
         * Properly authenticated: capture any pending message, then
         * release user data for next file.
         */
        msg = Iemail_grab_pending (udata);
        udata->active_file = NULL;
        up (&udata->sem);
    }

    /*
     * Release is atomic for this file. After returning from our
     * function, f will be added to the kernel's free list of files, but
     * we don't want pointers to our data floating around in the system.
     */
    f->private_data = NULL;
}

/* Deliver any pending message, ignoring failure. */
if (NULL != msg)
    (void)Iemail_deliver (msg, 0);

/*
 * Always succeeds; the return value is ignored by the kernel, anyway.
 * If we want to return errors, we need to use the flush method rather
 * than only catching release.
 */
return 0;
}

/* Iemail_fsync
 * DESCRIPTION: Implements the fsync and fdatsync system calls for Iemail.
 *              Extracts any pending message from the user and delivers it
 *              to its recipient.
 * INPUTS: f -- the file structure used for reading
 *          dent -- the directory entry
 *          datasync -- 1 for datasync call, 0 for sync call
 * OUTPUTS: none
 * RETURN VALUE: 0 on success
 *              -EPERM if user is not authenticated or has been deleted
 *                  after authentication
 *              -ENODATA if no message was being written
 *              -ENOENT if recipient of message does not exist
 * SIDE EFFECTS: removes message from user and delivers to recipient
 * SYNCHRONIZATION: Obtains the associated user's user data semaphore
 *                  in order to extract the message. Obtains both a
 *                  user list read lock and the recipient's semaphore
 *                  in order to deliver the message.
 */
static int
Iemail_fsync (struct file* f, struct dentry* dent, int datasync)
{
    Iemail_user_data_t* udata; /* user data structure */
    Iemail_msg_t* msg; /* message pending from user */

    /* If user has not authenticated yet, permission is denied. */
    if (NULL == (udata = f->private_data))
        return -EPERM;

    /* Capture any pending message for delivery. */
    down (&udata->sem);
    if (NULL == udata->active_file) {
        /* User was deleted after authentication. */
        up (&udata->sem);
        return -EPERM;
    }
    msg = Iemail_grab_pending (udata);
    up (&udata->sem);

    /* Try to deliver; return failure or success to caller. */
    if (NULL != msg)
        return Iemail_deliver (msg, 0);

    /* No message was being written! */
    return -ENODATA;
}

/*-----
 * The remainder of the file is module linkage.
 */

```

Iemail.c

```
/*
 * Use GCC's tagged field initialization to select only those functions
 * supported.
 */
static struct file_operations Iemail_fops = {
    owner:    THIS_MODULE,
    read:     Iemail_read,
    write:    Iemail_write,
    poll:     Iemail_poll,
    ioctl:    Iemail_ioctl,
    release:  Iemail_release,
    fsync:    Iemail_fsync,
};

module_init (Iemail_init);
module_exit (Iemail_exit);
EXPORT_NO_SYMBOLS;

MODULE_AUTHOR ("Steven S. Lumetta");
MODULE_LICENSE ("see file header");
```