# ECE391 Computer System Engineering Lecture 6

University of Illinois at Urbana- Champaign

Fall 2022

### Lecture Topics

- Shared resources
- Critical sections

#### MP1 Handin and Demo Schedule

- Code must be committed to master/main branch on GitLab by
  - 5:59 PM on Monday 9/12

#### Shared Data and Resources (1)

- The question
  - interrupt handlers and programs share resources
    - What resources are shared between them?
    - How might interactions cause problems?
    - What can we do to fix those problems?

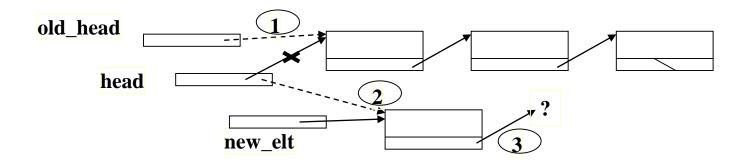
#### Thought Problem on Shared Resources (2)

- Obvious things
  - registers
    - solution? save them to the stack
  - memory
    - solution? privatize
    - will still need to share some things; discussed later

#### Thought Problem on Shared Resources (3)

- Less obvious
  - condition codes
    - solution? again, save them to the stack
  - shared data
- More subtle
  - external state (e.g., on devices)
  - compiler optimization (e.g., volatility)
  - security leaks
    - e.g., application waits for interrupt, then observes values written by OS to stack
    - solution? use separate stack for kernel

#### Example #1: a shared linked list



```
step 1: old_head = head;
```

Oops! an interrupt!

# Examples of Shared Resources: Example #1: a shared linked list

#### The problem?

- linked list structure has invariant
- head points to first, chained through last via next field, ends with NULL
- complete operation maintains invariant
- partial operation does not need atomic update

## Examples of Shared Resources: Example #2: external state

#### The core problem

- devices have state
- processors interact with devices using specific protocols
- protocol often requires several steps (e.g., I/O instructions)
- device cannot differentiate which piece of code performed an operation

#### Example:

- VGA controller operations for scrolling window with color modulation
- interrupt handler drives color manipulations
- program handles scrolling using pixel shift
- both use VGA attribute register (port 0x3C0)

## Examples of Shared Resources: Example #2: external state

- Protocol for attribute control register
  - 22 different attributes accessed via this register
  - first send index
  - then send data
  - VGA tracks whether next byte sent is index or data

- Problem: processor can't know which one is expected
- Solution: reading from port 0x3DA forces VGA to expect index next

#### Example #2: external state

- Consider the program code
  - the horizontal pixel panning register is register 0x13
  - assume that the code should write the value 0x03 to it

(discard)  $\leftarrow$  P[0x3DA] MOVW \$0x3DA, %DX

(%DX),%AL INB

 $0x13 \rightarrow P[0x3C0]$ MOVW \$0x3C0, %DX

MOVB \$0x13, %AL

OUTB %AL, (%DX)

 $0x03 \rightarrow P[0x3C0]$ MOVB \$0x03, %AL

OUTB %AL, (%DX)

## Examples of Shared Resources: Example #2: external state

- What happens if the interrupt occurs after the first write to 0x3C0?
  - the interrupt handler is executing basically the same code
  - leaves the VGA expecting an index

What is the solution?

#### Example #3: handshake synchronization

- A device generates an interrupt after it finishes executing a command
- Consider the following attempt to synchronize

```
the shared variable...
```

```
int device_is_busy = 0;
```

the interrupt handler...

```
device is busy = 0;
```

## Examples of Shared Resources: Example #3: handshake synchronization

The program function used to send a command to the device...

```
while (device_is_busy);/* wait until device is
    free */
device_is_busy = 1;
/* send new command to device */
```

- Q: Does the loop work?
- No.
  - Compiler assumes sequential program.
  - Variables can't change without code that changes them.

#### Example #3: handshake synchronization

```
LOOP: MOVL device_is_busy, %EAX

CMPL $0x0, %EAX

JNE LOOP
```

 Nothing can change variable, so no need to reload (move LOOP down a line).

#### Example #3: handshake synchronization

MOVL device is busy, %EAX

LOOP: CMPL \$0x0, %EAX

JNE LOOP

 Now nothing can change EAX, so move it down another line (to branch!).

#### Example #3: handshake synchronization

MOVL device is busy, %EAX

CMPL \$0x0, %EAX

LOOP: JNE LOOP

Will interrupt handler break you out of the resulting infinite loop?

# Examples of Shared Resources: Example #3: handshake synchronization

- Solution
  - mark variable as volatile
  - tells compiler to never assume that it hasn't changed between uses

the shared variable...

```
volatile int device is busy = 0;
```

- Why not mark everything volatile?
  - forces compiler to always re-load variables
  - more memory operations = slower program

# Examples of Shared Resources: Example #3: handshake synchronization

- Is it ok to swap setting the variable and sending the command?
- No.
  - introduces a race condition:

```
/* send new command to device */
---- INTERRUPT OCCURS HERE ----
device_is_busy = 1;
```

Next command call blocks (forever) for device to be free.

#### Critical Sections

- Some parts of program need to appear to execute atomically, i.e., without interruption
- Full version: atomic with respect to code in interrupt handler
  - for now, the clause is implied i.e., only interrupt handlers can operate during our programs
  - however, multiprocessors may have >1 program executing at same time

#### Critical Sections

- Solution?
  - IF (the interrupt enable flag)
     critical section start (CLI)
     (the code to be executed atomically)

- What else must be prevented?
  - no moving memory ops into or out of critical section!

critical section end (STI)

## Critical Sections in Examples

Example #2: external state

MOVW \$0x3C0, %DX the critical section should be as short as possible

OUTB %AL, (%DX)

MOVB \$0x03, %AL

OUTB %AL, (%DX)

## Critical Sections in Examples

- Why should critical sections be short?
  - avoid delaying device service by interrupt handler
  - long delays can even crash system (e.g., swap disk driver timeout)
- Example #1: a shared linked list

```
old_head = head;
head = new_elt;
new_elt->next = old_head;
could skip first statement,
but including is safer
STI
```

- If interrupt handler can change list, too, leaving out first inst. creates race
- Example #3: handshake synchronization—volatile suffices for this example