

ECE391

Computer System Engineering

Lecture 7



Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2022


Lecture Topics



- Multiprocessors and locks
- Spin locks

Multiprocessors and Locks

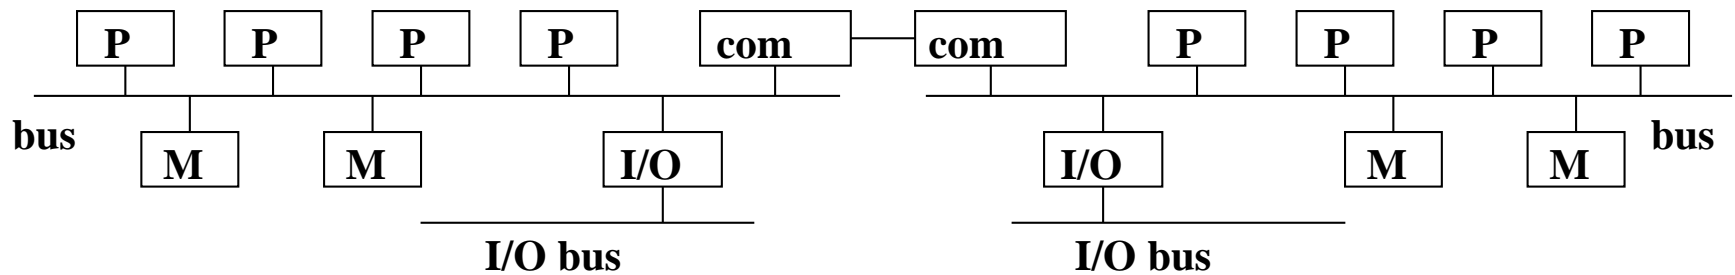
- We solved the critical section problem for uniprocessors
- What about multiprocessors?
 - *CLI ... critical section STI*
- What is a multiprocessor?
 - usually a symmetric multiprocessor (SMP)



- symmetric aspect: all processors have equal latency to all memory banks
- multicore processors are similar from our perspective

Multiprocessors and Locks (cont.)

- Some non-uniform memory architecture (NUMA) machines were built




Multiprocessors and Locks (cont.)


- Multithreaded code is not protected by *IF*
- Why haven't we solved the atomicity problem on multiprocessors?
 - interrupts **are masked** if *IF* is cleared!
 - answer: *IF* is **not cleared** on other processors!
 - just tell other processors to clear IF, too?
 - too slow
 - requires an interrupt!
- We need to use shared memory to synchronize...

Multiprocessors and Locks (cont.)

- Logically, we use a lock
 - when we want to access a piece of shared data we first look at the lock
 - if it's locked, we wait until it's unlocked




- once it's unlocked
 - we lock it
 - access the data
 - then unlock it



Multiprocessors and Locks (cont.)

- Locking must be atomic with respect to other processors!



INCL (%EAX)



read (%EAX)

add 1

put it back

LOCK: Prefix - execute following instruction with bus locked

LOCK INCL (%EAX)

An Example of Locking Implementation

/* The caller defines `int lock`. Calling `TestAndSet (&lock)` sets `lock` to 1 and returns the old value of `lock`. If `lock` is 0 then `TestAndSet(&lock)` returns 0 and sets the `lock` to 1. If the `lock` is already set to 1 by another process or thread, then 1 is returned. The caller keeps retrying `TestAndSet(&lock)` until it returns 0. */

TestAndSet:

.....

```
pushl %ebx
```

```
movl 8(%ebp), %ebx      # &lock to ebx
```

```
movl $1, %eax           # 1 (true) to eax
```

```
# Swap eax and lock, value 1(true) is copied to lock, eax receives old lock value
```

```
xchg %eax, (%ebx)       #Atomically exchange eax with lock.
```

```
#The atomicity of xchg guarantees that at most
```

```
#one thread holds the lock at any point in time
```

```
popl %ebx
```

.....

```
ret                     #return value (old value of lock) is already in eax
```

xchg => exchanges the contents of a register with the contents of any other register or memory location

Spin Locks

- The simplest lock

- spin lock
- lock op

do {

try to change lock variable from 0 to 1

} while (attempt failed);

- other work to do? ignore it!
- spin in a tight loop on the lock (hence the name)

- Once successful, program/interrupt handler **owns** the lock

Spin Locks (cont.)

- Only the owner can unlock
 - How?
 - (change lock variable to 0)
 - Need to be atomic?
 - (no, only owner can change when locked)

Linux Spin Lock API

- Static initialization

```
static spinlock_t a_lock = SPIN_LOCK_UNLOCKED;
```

- Dynamic initialization

```
spin_lock_init (&a_lock);
```

- When is dynamic initialization safe?
 - lock must not be in use (race condition!)
 - other synchronization method must prevent use

Linux Spin Lock API – Basic Functions



```
void spin_lock (spinlock_t* lock);
```

```
void spin_unlock (spinlock_t* lock);
```

Linux Spin Lock API – Testing Functions

```
int spin_is_locked (spinlock_t* lock) ;
```

returns 1 if held, 0 if not, but beware of races!

```
int spin_trylock (spinlock_t* lock) ;
```


make one attempt; returns 1 on success, 0 on failure

```
void spin_unlock_wait (spinlock_t* lock) ;
```

wait until available (race condition again!)

Linux Spin Lock API (cont.)

- Is spinlock enough to protect a critical section?



Linux Spin Lock API (cont.)

- What about ?




- Still need CLI/STI
- Which is first, CLI or lock?
 - CLI first
 - interrupt may occur between them, leading to scenario above

Linux' Lock/CLI Combo

```
static spinlock_t the_lock = SPIN_LOCK_UNLOCKED;  
unsigned long flags;
```

```
spin_lock_irqsave (&the_lock, flags);  
/* the critical section */  
spin_unlock_irqrestore (&the_lock, flags);
```



```
asm volatile ("  
    PUSHFL  
    POPL %0  
    CLI  
" : "=g" (flags) /* outputs */  
 : /* inputs */  
 : "memory" /* clobbers */  
);  
spin_lock (&the_lock);
```


Linux' Lock/CLI Combo (cont)

```
spin_unlock (&the_lock);  
asm volatile ("  
    PUSHL %0  
    POPFL  
  
    :                               /* outputs */  
    : "g" (flags)                  /* inputs  */  
    : "memory", "cc"               /* clobbers */  
);
```

Comments on code

- Notice that *spin_lock_irqsave* changes the flags argument (it's a macro)
- The “memory” argument
 - tells compiler that all memory is written by assembly block
 - prevents compiler from moving memory ops across assembly
- The “cc” argument
 - condition codes change
 - can lead to subtle bugs if left out!
- *spin_lock* and *spin_unlock* calls
 - become NOPs on uniprocessors
 - in that case, calls just change IF
- Restore rationale: may have had IF=0 on entry; if so, STI is unsafe

