ECE391 Computer System Engineering Lecture 10

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2022

Lecture Topics

- Programmable interrupt controller (PIC)
 - Linux 8259A Initialization
- Linux abstraction of PIC
- General interrupt abstractions
 - Interrupt Chaining
 - Soft Interrupts

ECE391 EXAM 1

- EXAM 1 September 27 (Tuesday);
 - Time: 7:00pm to 9:00pm
- Conflict Exam
 - Deadline to request conflict exam: Thursday, September 22 (by email to: <u>kalbarcz@Illinois.edu</u>)

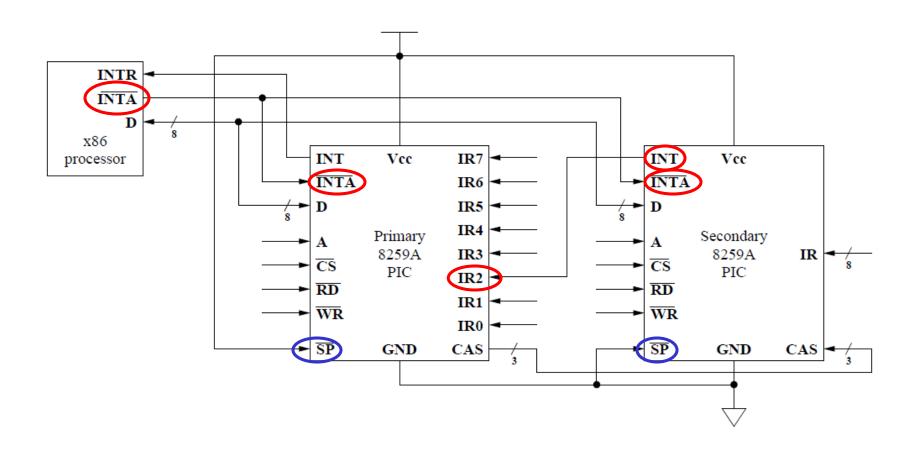
- Exam 1 Review Session in collaboration with HKN
 - <u>Date:</u> Saturday September 24;
 - <u>Time:</u> from 2:30 to 4:30pm
 - Location: ECEB 1002

ECE391 EXAM 1

- Topics covered by EXAM 1
 - Material covered in lectures (Lecture1 Lecture10)
 - x86 Assembly
 - C-Calling Convention
 - Synchronization
 - Interrupt control (using PIC)
 - Material covered in discussions
 - MP1

NO Lecture on Tuesday, September 27

Cascade Configuration of PICs



PIC (cont.)

- Previous figure showed x86 configuration of two 8259A's
 - primary 8259A mapped to ports 0x20 & 0x21
 - secondary 8259A mapped to ports 0xA0 & 0xA1
 - secondary PIC connects to IR2 on primary PIC

Question

- prioritization on 8259A: 0 is high, 7 is low
- what is prioritization across all 15 pins in x86 layout?
- (highest) P0...P1...S0...S7...P3...P7 (lowest)

PIC (cont.)

- In Linux (initialization code to be seen shortly)
 - primary PIC IR's mapped to vector #'s 0x20 0x27
 - secondary PIC IR's mapped to vector #'s 0x28 0x2F
 - remember the IDT?

	0x00	division error	
		division error	
	:		
	0x02	NMI (non-maskable interrupt)	
	0x03	breakpoint (used by KGDB)	
0x00-0x1F	0x04	overflow	
	:		
defined	0x0B	segment not present	
by Intel	0x0C	stack segment fault	
	0x0D	general protection fault	
	0x0E	page fault	
	:		
	0x20	IRQ0 — timer chip	
	0x21	IRQ1 — keyboard	
0x20-0x27	0x22	IRQ2 — (cascade to secondary)	
	0x23	IRQ3	
primary	0x24	IRQ4 — serial port (KGDB)	
8259 PIC	0x25	IRQ5	
	0x26	IRQ6	example
	0x27	IRQ7	of
	0x28	IRQ8 — real time clock	possible
	0x29	IRQ9	settings
0x28-0x2F	0x2A	IRQ10	
	0x2B	IRQ11 — eth0 (network)	
secondary	0x2C	IRQ12 — PS/2 mouse	
8259 PIC	0x2D	IRQ13	
	0x2E	IRQ14 — ide0 (hard drive)	
	0x2F	IRQ15	
0x30-0x7F	÷	APIC vectors available to device drivers	
0x80	0x80	system call vector (INT 0x80)	
0x81-0xEE	:	more APIC vectors available to device drivers	
0xEF	0xEF	local APIC timer	
0xF0-0xFF	:	symmetric multiprocessor (SMP) communication vectors	

Interrupt Descriptor Table (IDT)

Linux 8259A Initialization

```
void init 8259A(int auto eoi)
        unsigned long flags;
        i8259A auto eoi = auto eoi;
        spin lock irqsave(&i8259A lock, flags)
        outb(0xff, 0x21);
                                /* mask all of 8259A-1 */
        outb(0xff, 0xA1);
                                 /* mask all of 8259A-2 */
         * outb p - this has to work on a wide range of PC hardware.
        outb p(0x11, 0x20);
                                /* ICW1: select 8259A-1 init */
        outb p(0x20 + 0, 0x21); /* ICW2: 8259A-1 IRO-7 mapped to 0x20-0x27 */
        outb p(0x04, 0x21);
                                /* 8259A-1 (the primary) has a secondary on IRQ2
        if (auto eoi)
                outb p(0x03, 0x21);
        else
                outb p(0x01, 0x21);
        outb p(0x11, 0xA0);
                                /* ICW1: select 8259A-2 init */
        outb p(0x20 + 8, 0xA1); /* ICW2: 8259A-2 IR0-7 mapped to 0x28-0x2f */
        outb_p(0x02, 0xA1);
                                / 8259A-2 is a secondary on primary's IRQ2
        outb p(0x01, 0xA1);
        if (auto eoi)
                 * in AEOI mode we just have to mask the interrupt
                 * when acking.
                i8259A irq type.ack = disable 8259A irq;
        else
                i8259A irq type.ack = mask and ack 8259A;
                                /* wait for 8259A to initialize */
        udelay(100);
        outb(cached 21, 0x21);
        outb(cached A1, 0xA1);
        spin unlock irqrestore(&i8259A lock, flags)
```

Comments on Linux' 8259A Initialization Code

- 1. What is the auto_eoi parameter? always = 0
- 2. Four initialization control words to set up primary 8259A
- 3. Four initialization control words to set up secondary 8259A

```
port(A=?) info contained in Initialization Control Word

ICW1 0 start init, edge-triggered inputs, cascade mode, 4 ICWs

ICW2 1 high bits of vector #

ICW3 1 primary PIC: bit vector of secondary PIC;

secondary PIC: input pin on primary PIC

ICW4 1 ISA=x86, normal/auto EOI
```

Comments on Linux' 8259A Initialization Code (cont.)

- 5. What does the "_p" mean on the "outb" macros?
 - add PAUSE instruction after OUTB; "REP NOP" prior to P4
 - delay needed for old devices that cannot handle processor's output rate
- 6. Critical section spans the whole function
 - avoid other 8259A interactions during initialization sequence
 - (device protocol requires that four words be sent in order)
- 7. Why use _irqsave for critical section?
 - this code called from other interrupt initialization routines
 - which may or may not have cleared IF on processor

Linux Abstraction of PICs

- Uses a jump table
 - same as vector table (array of function pointers)

- Table is <u>hw_irq_controller</u> structure (or struct irq_chip)
 - each vector # associated with a table
 - table used to interact with appropriate PIC (e.g., 8259A, or Advanced PIC)

human-readable name
startup function
shutdown function
enable function
disable function
mask function
mask_ack function
unmask function
(+ several others)

Linux Abstraction of PICs

- hw_irq_controller structure definition
 - IRQs are #'d 0-15 (correspond to vector # 0x20)

```
const char* name;
unsigned int (*startup) (unsigned int irq);
void (*shutdown) (unsigned int irq);
void (*enable)...

void (*disable)...

void (*ack)...

void (*end)...
/* we'll ignore the others... */
```

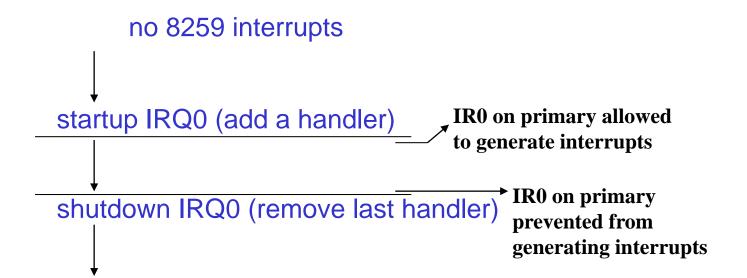
PIC Functions in Jump Table: Explanation

 Initially, all 8259A interrupts are masked out using mask on 8259A

- startup and shutdown functions
 - startup is called when first handler is installed for an interrupt
 - shutdown is called after last handler is removed for an interrupt
 - both functions change the corresponding mask bit in 8259A implementation

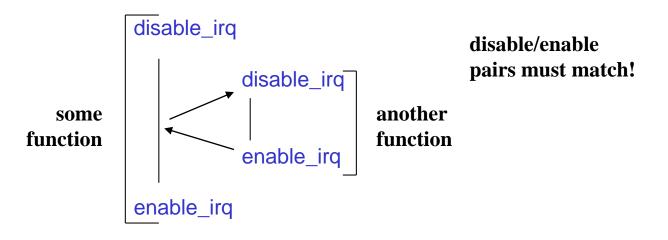
PIC Functions in Jump Table: Explanation (cont.)

Example



PIC Functions in Jump Table (cont.)

- disable/enable functions
 - used to support nestable interrupt masking (disable_irq, enable_irq)



- on 8259
 - first disable_irq calls jump table disable, which masks interrupt on PIC
 - last enable_irq calls jump table enable, which unmasks interrupt on PIC

PIC Functions in Jump Table (cont.)

ack function

- called at start of interrupt handling to ack receipt of the interrupt
- on 8259 (mask and ack), masks interrupt on PIC, then sends EOI to PIC

end function

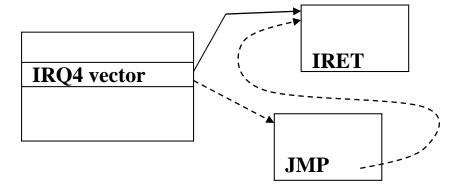
- called at end of interrupt handling
- on 8259, enables interrupt (unmasks it) on PIC

General Interrupt Abstractions: Interrupt Chaining

- Hardware view: 1 interrupt → 1 handler
- Problems
 - may have > 15 devices
 - > 1 software routines may want to act in response to device
 - examples:
 - hotkeys for various functions
 - move mouse to lower-right corner to start screen-saver

General Interrupt Abstractions: Interrupt Chaining (cont.)

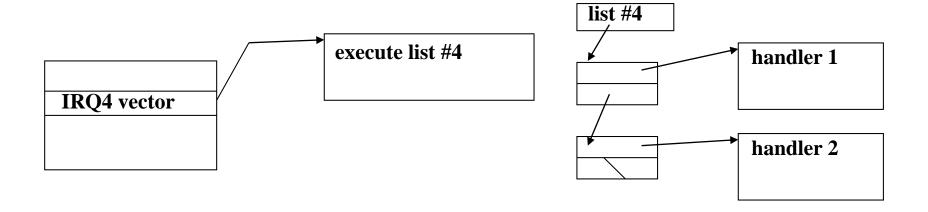
- One approach
 - used by terminate and stay resident (TSR) programs in DOS
 - form linked list (chain) of handlers using JMP instructions
 - not very clean
 - no way to remove self
 - unless you're first in list
 - to be fair
 - TSR program not designed for removal



General Interrupt Abstractions Interrupt Chaining (cont.)

Solution

- interrupt chaining with linked list data structure
- (not list embedded into code!)



General Interrupt Abstractions: Interrupt Chaining (cont.)

- Drawbacks of chaining
 - for > 1 device
 - must query devices to see if they raised interrupt
 - not always possible
 - for 1 device
 - must avoid stealing data/confusing device
 - example
 - by sending two characters to serial port
 - in response to interrupt declaring port ready for one char.

General Interrupt Abstractions: Soft Interrupts (cont.)

- Recall: why support interrupts?
 - slow device gets timely attention from fast processor
 - processor gets device responses without repeatedly asking for them
- A useful concept in software
 - example: network encryption/decryption
 - packet arrives, given to decrypter
 - when decrypter (software program) is done
 - want to interrupt program
 - to transfer data from packet
 - but has no access to INTR pin

General Interrupt Abstractions: Soft Interrupts (cont.)

Solution

- software-generated (soft) interrupt
- (similarly, but later, signals—user-level soft interrupts)
- runs at priority between program and hard interrupts
- usually generated
 - by hard interrupt handlers
 - to do work not involving device

Linux version is called tasklets

- used by code provided to you for MP1
- discussed later