ECE391 Computer System Engineering Lecture 12

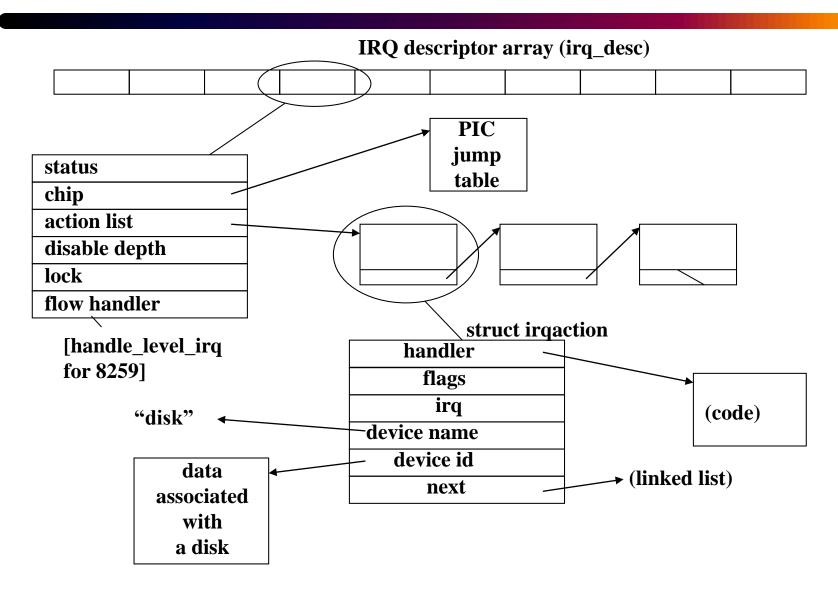
Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Spring 2021

Lecture Topics

- Linux interrupt system
 - data structures
 - handler installation & removal
 - invocation
 - execution

Linux' Interrupt Data Structures



```
I inux' request_irq
int
                                            irg # (0-15 for PIC)
request irq (unsigned int irq,
                                                                     (kernel/irq/manage.c)
                                              pointer to interrupt handler
              irq handler t handler,
              unsigned long irqflags,
                                              bit vector
              const char* devname,
                                              human-readable device name
              void* dev id)
                                              pointer to device-specific
                                              data
    struct irgaction* action;
    int retval;
    if (irq >= NR IRQS)
        return -EINVAL;
    if (!handler)
        return -EINVAL;
    action = kmalloc (sizeof (struct irqaction), GFP ATOMIC); - Dynamically allocate
    if (!action)
                                                                             new action structure for
        return -ENOMEM;
                                                                             linked list of actions
                                                     handler
    action->handler = handler;
                                                      flags
                                                                 Fill in the new
    action->flags
                      = irqflags;
                                                       irq
    cpus clear (action->mask);
                                                                 action structure
                                                  device name
    action->name
                      = devname;
                                                    device id
    action->next
                      = NULL;
                                                      next
    action->dev id
                      = dev id;
    select smp affinity (irq);
                                                   Add the new action structure
    retval = setup irq (irq, action)
                                                   to the link list of actions
    if (retval)
        kfree (action);
                                                   Upon failure free the action structure
    return retval;
```

Comments on request_irq - Arguments

- irq irq # (0-15 for PIC)
- handler pointer to interrupt handler with following arguments
 - irq #
 - dev_id pointer (type void*)
 - handlers should return 1 if handled, 0 if not
- irqflags bit vector
 - IRQF_SHARED interrupt chaining is allowed
 - IRQF_DISABLED execute handlers with IF=0
 - IRQF_SAMPLE_RANDOM use device timing as source of random numbers
- devname human-readable device name (visible in /proc/interrupts)
- dev id pointer to device-specific data (returned to handler when called)

Comments on request_irq (cont.)

- Start by checking values (two checks from the shown code)
- Dynamically allocate new action structure for linked list of actions
- Fill in the new action structure
- Try to add it
 - using setup_irq
 - if call fails, free the action structure

```
int
setup irq (unsigned int irq, struct irqaction* new)
    struct irq desc* desc = irq desc + irq;
                                                                                    Linux setup_irq
    struct irgaction* old;
    struct irgaction** p;
    unsigned long flags;
    int shared = 0;
                                                                                               (kernel/irq/manage.c)
    if (irq >= NR IRQS)
        return -EINVAL;
    if (desc->chip == &no irq chip)
        return -ENOSYS;
    if (new->flags & IRQF SAMPLE RANDOM)
        rand_initialize_irq (irq);
                                                                      Critical section begins
    spin lock irgsave (&desc->lock, flags);
    p = &desc->action;
    if ((old = *p) != NULL) {;
        /* Can't share interrupts unless both agree to and are same type. */
        if (!((old->flags & new->flags) & IRQF SHARED) ||
            ((old->flags ^ new->flags) & IRQF_TRIGGER_MASK)) {
            spin unlock irgrestore (&desc->lock, flags);
            return -EBUSY;
        /* add new interrupt at end of irq queue */
        do {
                                                                       New action is added at the end of the link list
            p = &old->next;
            old = *p;
        } while (old);
                                                                           Interrupt
                                                                                         status
        shared = 1;
                                                                           descriptor
                                                                                         chip
                                                                                         action list
    *p = new;
                                                                                         disable depth
    if (!shared)
                                                                                                                           PIC
        irq chip set defaults (desc->chip);
                                                                                         lock
        desc->status &= "(IRQ AUTODETECT | IRQ WAITING | IRQ INPROGRESS);
                                                                                                                           jump table
                                                                                         flow handler
        if (!(desc->status & IRQ_NOAUTOEN)) {
            desc->depth = 0;
                                                                                                              human-readable name
            desc->status &= "IRO DISABLED:
                                                            Startup PIC for this interrupt
                                                                                                              startup function
            desc->chip->startup (irq);
                                                                                                              shutdown function
        } else
                                                                                                              enable function
            /* Undo nested disables: */
                                                                                                              disable function
            desc->depth = 1;
                                                                                                              mask function
                                                                                                              mask ack function
                                                                                                              unmask function
                                                                    Critical section ends
    spin unlock irgrestore (&desc->lock, flags);
                                                                                                              (+ several others...)
    new->irg = irg:
                                                             Create new /proc/irg/<irg#> directory
    register irg proc (irg);
    new->dir = NULL:
                                                             Add handler subdirectory in /proc/irq/<irq#>/<action name>
    register handler proc (irq, new);
    return 0;
```

Comments on stup_irq

Start with a couple of sanity checks

- If random sampling flag is set
 - initialize as source of random numbers
 - good random numbers are hard to find!
 - devices such as disks can be used because of "random" rotational latency to read data (for example)

Comments on stup_irq (cont.)

- Critical section (most of function)
 - blocks other activity on descriptor
 - uses irqsave/restore to allow handlers to be added from any context

 Note that new action goes at the <u>end</u> of the linked list, not the start

- Interrupt chaining only allowed
 - if <u>all</u> handlers agree to it
 - otherwise only first handler is successfully added

Comments on stup_irq (cont.)

- If this handler is the first
 - make sure that PIC jump table has proper default functions
 - clear some status flags
 - clear any previous software disablement (depth)
 - call the PIC startup function

- After critical section
 - create /proc/irq/<irq#> directory if necessary
 - add handler subdirectory in /proc/irq/<irq #>/<action name>

```
void
free irq (unsigned int irq, void* dev id)
                                                                               Linux free_irq
    struct irq desc* desc;
    struct irqaction** p;
                                                                                  (kernel/irq/manage.c)
    unsigned long flags;
    if (irq >= NR_IRQS)
        return;
                                                     Find the correct irq descriptor
    desc = irq desc + irq;
    spin lock irqsave (&desc->lock, flags);
                                                          Critical section begins
    p = &desc->action;
    for (;;) {
        struct irgaction* action = *p;
        if (action) {
                                                   Search for action in list
            struct irgaction** pp = p;
            p = &action->next;
                                                    The "continue" causes the rest of the statement
            if (action->dev id != dev id)
                                                    body in the loop to be skipped.
                continue:
            /* Found it - now remove it from the list of entries */
            *pp = action->next;
            if (!desc->action) {
                                                              No more handlers for the IRQ:
                desc->status |= IRO DISABLED;
                                                               you can invoke PIC shutdown
                desc->chip->shutdown (irq);
            spin unlock irqrestore (&desc->lock, flags);
                                                                      Critical section ends
            /* Remove from /proc/irg/<irg #> */
                                                            Remove handler subdirectory in /proc/irg/<irg#>/<action name>
            unregister handler proc (irq, action);
            /* Make sure it's not being used on another CPU */
            synchronize irq (irq);
                                Free the action structure
            kfree (action);
            return;
        printk (KERN ERR "Trying to free already-free IRQ %d\n", irq);
        spin unlock irqrestore (&desc->lock, flags);
                                                                             Critical section ends
        return;
```

Comments on free_irq

- Arguments
 - **irq irq** # (0-15 for PIC)
 - dev_id MUST match pointer value used in request_irq call

Start by checking arguments

- Critical section starts to prevent manipulations of descriptor
 - blocks interrupt from <u>starting</u> to execute
 - (interrupt may <u>already</u> be executing on another processor)

Comments on free_irq (cont.)

- Search for action in list
 - based on dev id pointer
 - hence need for matching pointer
 - passing NULL dev_id not allowed with chained interrupts
- Once found (notice that most of the loop executes exactly once)
 - remove the action
 - if this action/handler was the last for this interrupt,
 - turn on software disablement
 - (doing so signals interrupts waiting for descriptor lock to abort once they obtain the lock)
 - call the PIC shutdown function

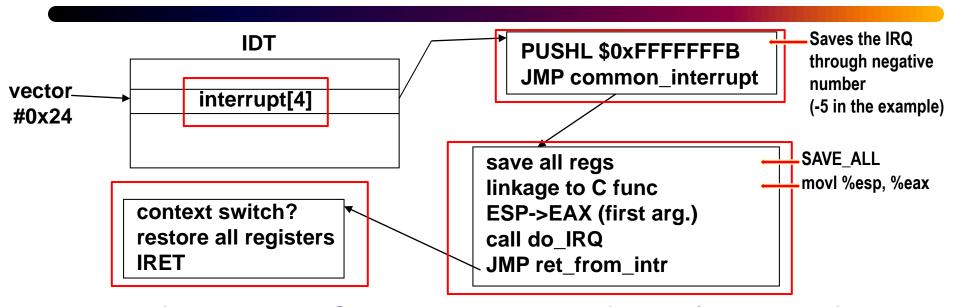
Comments on free_irq (cont.)

• Remove the /proc/irq/<irq #>/<action name> entry

- On an SMP
 - interrupt may be executing on another processor
 - need to wait for it to finish after releasing descriptor lock

Finally, free the action structure

Interrupt Invocation



- Why funnel all IRQs into one piece of code(instead of producing one function per IRQ)?
 - kernel code/size versus speed tradeoff
 - perhaps no big deal for 16
 - keep in mind that you're saving a tiny number of instructions
 - APIC used by most SMPs has 256 vectors

	0x00	division error	
	:		
	0x02	• • • • • • • • • • • • • • • • • • • •	
	0x03		
0x00-0x1F	0x04	overflow	
	:		
defined	0x0B	segment not present	
by Intel	0x0C	stack segment fault	
		general protection fault	
	0x0E	page fault	
	:		
	0x20	IRQ0 — timer chip	
	0x21	IRQ1 — keyboard	
0x20-0x27	0x22	IRQ2 — (cascade to slave)	
	0x23	IRQ3	
master	0x24	IRQ4 — serial port (KGDB)	
8259 PIC	0x25	IRQ5	
		IRQ6	example
		IRQ7	of
		IRQ8 — real time clock	possible
	l .	IRQ9	settings
0x28-0x2F		IRQ10	
		IRQ11 — eth0 (network)	
slave		IRQ12 — PS/2 mouse	
8259 PIC		IRQ13	
		IRQ14 — ide0 (hard drive)	
	0x2F	IRQ15	
0x30-0x7F	:	APIC vectors available to device drivers	
0x80	0x80	system call vector (INT 0x80)	
0x81-0xEE	:	more APIC vectors available to device drivers	
0xEF	0xEF	local APIC timer	
0xF0-0xFF	:	symmetric multiprocessor (SMP) communication vectors	

Interrupt Descriptor Table

