ECE391 Computer System Engineering Lecture 14

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Spring 2021

Lecture Topics

- Virtual memory motivation
- x86 support for VM
 - protection model
 - segmentation
 - paging

Aministrivia

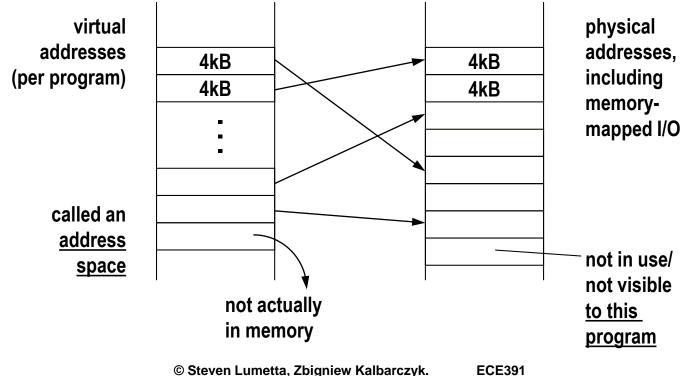
- MP2 Checkpoint 2
 - Due by 5:59pm Monday, March 15

MP3 Teams

- Please submit a google form: https://forms.gle/FfeSuDFjo49VNjtd7 before Thursday 3/11 at 11:59PM CT
- Every team must have 4 members

Virtual Memory Definition

- What is virtual memory?
 - <u>indirection</u> between memory addresses seen by software and those used by hardware
 - typically done with large blocks, e.g., 4kB or 4MB in x86

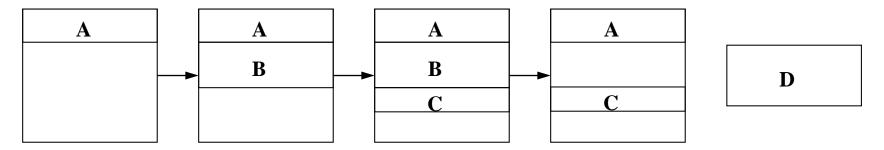


Virtual Memory Motivation

- Why use virtual memory? (in decreasing order of importance)
 - protection
 - one program <u>cannot</u> accidentally or deliberately destroy another's data
 - the memory is simply not accessible
 - more effective sharing
 - two (or more) programs that share library code can share a single copy of the code in physical memory
 - code and data not actively used by a program
 - can be pushed out to disk
 - to make room for other programs' active data
 - provides the illusion of a much larger physical memory

Virtual Memory Motivation (cont.)

- no fragmentation (little to none)
 - systems without virtual memory suffer fragmentation when try to multitask
 - e.g., if we run A followed by B followed by C, and B finishes, we can't give
 D a contiguous block of memory, even though it fits in the absolute sense



simplifies program loading and execution:
 no relocation of code, rewriting stored pointer values, etc.

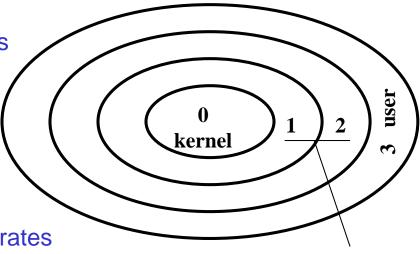
x86 Support for VM

- protection model
- segmentation
- paging



x86 Protection Model

- Four rings: kernel (ring 0) through user (ring 3)
 - lower numbers are more privileged
 - lower numbers <u>never</u> call/trust higher numbers
 - higher numbers call lower numbers only through narrow interfaces (e.g., system calls)
- CPL current privilege level (of executing code)
- RPL requestor's privilege level; when code at high privilege level executes on <u>behalf</u> of code at lower level, some accesses may voluntarily lower privilege to that of caller/beneficiary
- DPL descriptor privilege level;
 level necessary to execute code/data
- if MAX(CPL,RPL) > DPL, processor generates an exception (general protection fault)



not used by Linux

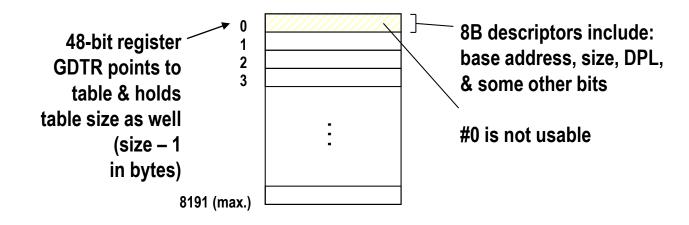
x86 Segmentation

- x86 actually has two levels of indirection
- One is mostly unused...(this one!) => Segmentation
- A segment is a contiguous portion of a linear address space
 - such as the 32-bit space of physical addresses
- x86 in protected mode <u>always</u> uses segmentation

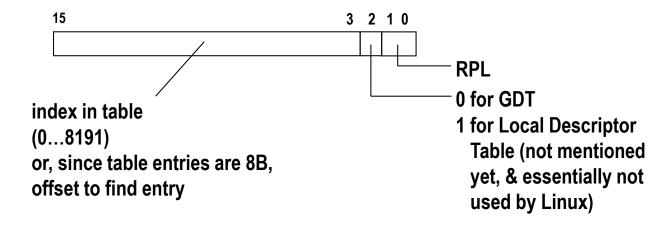
code segment	CS	each segment register has 16 bits visible + ~64 bits shadow (not accessible via ISA) that cache the description of the segment # referenced by the visible 16 bits
data segment	DS	
extra data segment	ES	
still more extras (floating point + another)	FS	
	GS	
stack segment	SS	

x86 Segmentation (cont.)


Global descriptor table (GDT)



Segment Register Meaning

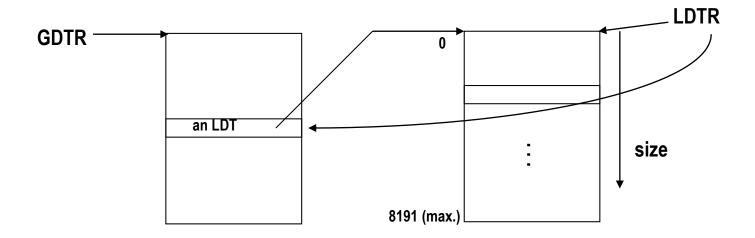


Segment Register Meaning (cont.)



Segment Register Meaning (cont.)

- GDT entries can also describe local descriptor tables (LDTs)
 - LDT originally meant to be per-task segment tables
 - LDTR points to current LDT (includes base, size, and index of LDT in GDT)



Segment Register Meaning (cont.)

- Descriptors can also differentiate
 - code (executable and possible readable) from data (readable and possibly writable)
 - other useful things

 Also descriptors in the GDT can describe certain aspects of program state (e.g., the task state segment, or TSS)

Linux Use of Segmentation (details in asm/segment.h)

