ECE391 Computer System Engineering Lecture 19

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2022

Lecture Topics

System Calls

ECE391 EXAM 2

- EXAM 2 November 1 (Tuesday)
 - Time: 7:00pm to 9:00pm
- Conflict Exam

– Deadline to request conflict exam:

Thursday, October 27, by 5:00pm

(by email to: <u>kalbarcz@Illinois.edu</u>)

ECE391 EXAM 2

- Exam 2 Review Session in Class:
 - During the Lecture Time on Thursday, October 27
- Exam 2 Review Session in collaboration with HKN
 - TBD

ECE391 EXAM 2: Topics

Topics covered by the EXAM 2

- Material covered in lectures
 (Lecture 12, Lecture 13, Lecture 14, Lecture 15, Lecture 16, Lecture 17, Lecture 19)
- MP2 (ModeX, Tux controller)
- MP3.1 and MP3.2 (file system)
- Material covered in discussions

NOTE: Scheduling NOT included in Exam 2

NO Lecture on Tuesday, November 1

System Calls

 Unix/Linux uses system calls to implement most interfaces between User Mode processes and hardware device

- System calls need protection boundary (user to kernel mode)
 - supported by ISA
 - need assembly linkage

System Call

- In Linux, all system calls use the following conventions
 - INT \$0x80 # to invoke (IDT vector 0x80)
 - EAX = system call # (asm/unistd.h)
 - EAX = return value (negative for errors) (asm-generic/errno.h)

System Call

IRQ8 - real time clock

IRQ11 - eth0 (network)

IRQ14 - ide0 (hard drive)

APIC vectors available to device drivers system call vector (INT 0x80)

IRO12 - PS/2 mouse

IRQ9

IRO10

IRO13

0x2E

Vector in IDT is system_call (in arch/i386/kernel/entry.S)

IDT

Table

- saves registers to stack
- check for a valid system call #
- call specific routine using a jump table:

stack (in kernel level) now appears as shown...(note the argument order)

- on completion of the system call handler
 - get the return code from eax
 - restore all the registers and IRET to resume the User Mode process

ret. addr
orig. EBX
orig. ECX
orig. EDX
orig. ESI
orig. EDI
orig. EBP

example

possible

settings

Example: From Lib Code to System Call

- C library code
 - arranges the arguments
 - before performing the system call
 - for example, open and sys_open

Open System Call Wrapper (1)

```
open: pushl
                                      # save EBX to stack collee saved register
                %ebx
      movl
                0x10(%esp),%edx
                                      # EDX mode
                0x0C(%esp),%ecx
                                      # ECX flags
      movl
                0x08(%esp),%ebx
                                      # EBX name
      movl
                $0x05,%eax
      movl
                                      # open is system call #5
      int
                $0x80
                                      # do the system call
                $0xFFFFF001,%eax
                                      # -1 to -4095 are errors
      cmpl
      jb
                done
                                      # others are valid descriptors
      xorl
                %edx,%edx
                                      # negate error number
      subl
                %eax,%edx
      pushl
                %edx
                                      # save EDX (caller-saved)
      call
                  errno location
                                      # get pointer to errno
      popl
                %ecx
                                      # pop error number into ECX
                %ecx,(%eax)
                                      # save error number in errno
      movl
      orl
                $0xFFFFFFFF,%eax
                                      # return -1
                                      # restore EBX from stack
done: popl
                %ebx
      ret
```

Open System Call Wrapper

 Call to open translates to a call to sys_open inside the kernel

- Before going on, let's see error handling path
- Recall that, on error, C library
 - returns -1 (for many calls, including open)
 - stores error code in errno variable

 But library code is relocatable, so errno address is not fixed

Open System Call Wrapper (2)

- Find static data (in this case errno) in a relocatable library
 - the library may be loaded to different parts of memory from the program's point of view
 - the address must be adjusted for the current execution

```
__errno_location:

call getIP # EAX <- raddr

raddr: addl $errno - raddr, %eax # adjust return value

ret

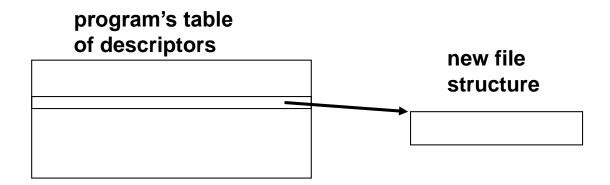
getIP: movl (%esp), %eax # read return address into eax

ret
```

- make a "fake" call
- call pushes return address onto stack
- return address is located at fixed offset from static variable
- load from stack and add offset to obtain pointer to variable, errno in our case

More on sys_open

- sys_open (in fs/open.c) invokes do_sys_open
- do_sys_open does the following
 - get a free file descriptor
 - open the file (using do_filp_open)
 - attach the file to the descriptor



Open System Call Wrapper

filp open

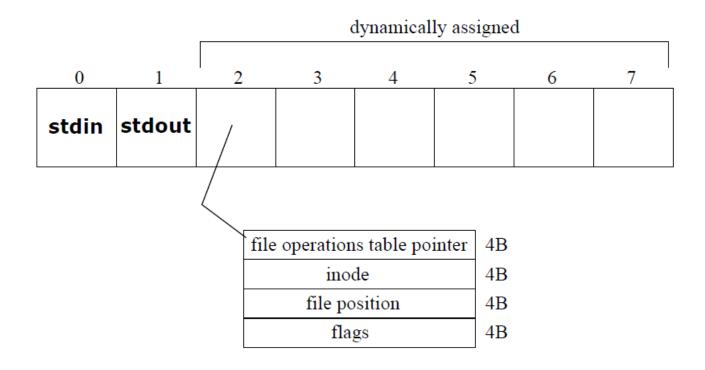
- check access rights
- find VFS mount
- get a file structure from a free list
- open directory entry (using <u>dentry</u> open, below)

__dentry_open

- fill in file structure
- if fops structure is not NULL && fops->open entry is not NULL,
 call fops->open

MP3 File System Abstractions

- Each task can have up to 8 open files
- Open files are represented with a file array (in a Process Control Block; PCB)
- File array is indexed by a file descriptor



MP3: Process Control Block (PCB)

- Process Control Block (PCB) is analogous to a task struct (process descriptor) in Linux
- PCB for each process is stored above the kernel stack corresponding to this process

