ECE391 Computer System Engineering Lecture 22

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign
Spring 2021

Lecture Topics

- Memory allocation interfaces in the kernel
 - kmalloc
 - slab caches
 - vmalloc
 - buddy system

Aministrivia

- MP3 Checkpoint 3
 - Due by 6:00pm Monday, April 12

Files for Memory Management

- headers (all under linux/): gfp.h, slab.h, vmalloc.h, slab_def.h, slub_def.h
- sources: mm/slab.c, mm/vmalloc.c
- swap-related: mm/swap.c, mm/swapfile.c, mm/page_alloc.c

Memory Management

- Paging translates virtual address to physical address
- Some portion of the physical memory permanently assigned to the kernel: stores code and data
- Rest of the physical memory can be allocated at runtime dynamic allocation

 Kernel must keep track of the status of each memory page in physical memory, e.g., is the page free

Memory Zones (1)

- In allocating memory Linux kernel must deal with two important hardware constrains of the x86 architecture:
 - Direct Memory Access (DMA): older processor can address only the first 16MB of RAM
 - In 32-bit computers with a lot of memory, the processor cannot directly access all physical memory because the liner address apace is too small

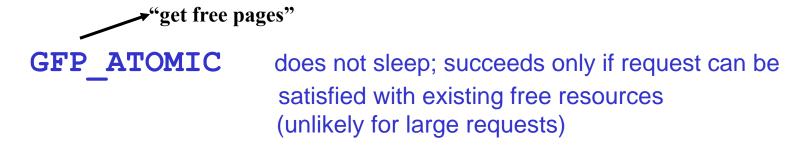
Memory Zones (2)

- To address hardware limitations Linux partitions the physical memory into three zones:
- ZONE_DMA
 - Contains memory pages below 16MB
- ZONE_NORMAL
 - Contains memory pages at and above 16MB and below 896MB
- ZONE_HIGHMEM
 - Contains memory pages at and above 896MB

Memory Allocation

Overview

- a few small items → kmalloc
- a lot of items, repeatedly → slab cache
- a big, physically contiguous region → free pages
- a big area of virtual memory → vmalloc (not necessarily physically contiguous)
- flags/allocation priorities (common to all interfaces)



Memory Allocation

may sleep to wait for pages

GFP KERNEL by kernel, drivers, etc.

GFP_NOFS no file system calls (avoids pushing pages to disk)

GFP NOIO no I/O operations at all

GFP USER on behalf of a user (low priority)

GFP DMA DMA accessible (low physical addresses

on some machines)

GFP HIGHMEM high memory (PAE (phys. address extensions) on

x86) is acceptable

(two underscores)

Basic Interface

```
void* kmalloc (size_t size, gfp_t flags);
```

- uses exponentially-sized slab caches (to be discussed)
 - ranging from 8B to several MB
 - up to 4MB in our kernel
- each allocation is contiguous in physical memory

Basic Interface

- Managing a private cache of objects (slab cache)
 - frequent allocations/deallocations
 - one cache per item type
 - physical memory is contiguous
 - protocol
 - Creation returns a page handle
 - Use handle to allocate/deallocate objects or to destroy the slab cache when done

Basic Interface

- name used to avoid >1 cache for same structure
- size is object size; cache grows/shrinks automatically
- alignment specified for individual objects

Flags for slab cache allocation

SLAB_HWCACHE_ALIGN

align objects to cache lines (makes accesses a little faster)

SLAB_CACHE_DMA use DMA-accessible memory

SLAB_POISON fill new memory with 0xA5A5A5A5

SLAB_RED_ZONE bound objects with "red zones" (test buffer overruns)

Slab cache constructor function

- Callback function used when memory is allocated for the slab cache
- called for each object in new slab
- NOT called for each object allocation (kmem_cache_alloc)
- third argument used to be flags (now always 0)

Slab cache API

Slab cache allocation/deallocation and destruction

- flags are passed to lower allocator (kmalloc) iff a new slab is allocated
- zalloc version zeroes memory in new object

```
void kmem cache free (kmem cache*, void*);
```

Slab cache API

```
void kmem_cache_destroy (kmem_cache*);
```

fails silently (logs error message) if not empty

```
int kmem_cache_shrink (kmem_cache*);
```

frees empty slabs; returns 0 if all slabs released

Getting big chunks of memory

- multiples of page size (4kB on x86; ISA-dependent)
- physically contiguous

- flags are same as for kmalloc
- order is log (base 2) of number of pages requested
- (the latter two function names start with two underscores)

Getting big chunks of memory

```
void free_page (unsigned long);
void free_pages (unsigned long, int order);
```

- order <u>must match</u> value used when allocated!
- These functions do <u>not</u> check for you!

Getting big chunks of memory

- Virtual memory allocation (request size in bytes, but allocates pages)
 - all functions return virtual addresses
 - but all other functions discussed today allocate physically contiguous regions
 - what if we don't care (or need a bigger region)?
 - use vmalloc (see linux/vmalloc.h)

```
void* vmalloc (unsigned long size);
void vfree (void* addr);
```

Memory Fragmentation

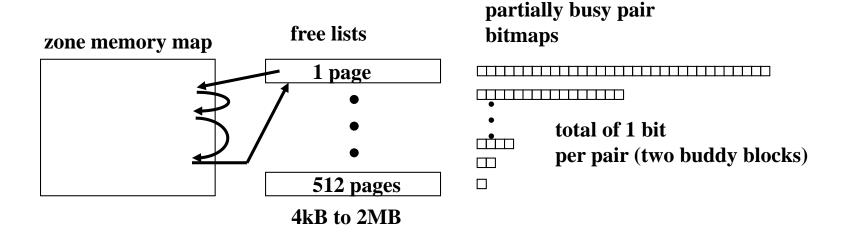
• External fragmentation: Frequent allocation/deallocation of group of continuous pages in physical memory of different sizes may lead to situation in which several small blocks of free pages are scattered inside blocks of allocated pages. As a result it may be impossible to allocate a large block of contiguous pages even if there is enough free pages.

 Internal fragmentation: caused by a mismatch between the size of the memory request and the size of the memory area allocated to satisfy the request

The Buddy System

- Problem: how to implement memory allocation inside the kernel
 - need page alignment for allocations
 - may need contiguous regions of physical memory
 - need flexible allocation granularity
 - want to avoid always rewriting page tables
 - can't easily add information to allocated (physical) memory chunks

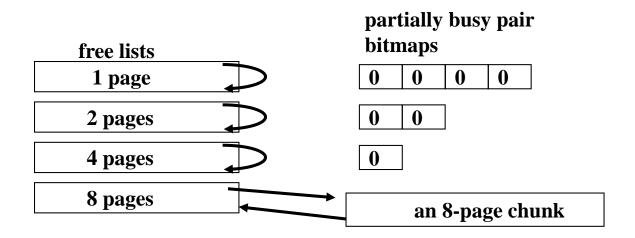
The Buddy System



The Buddy System

- Traditional simple answer
 - exponential bins: 1 page, 2 pages, 4 pages, etc.
 - buddy system extends with dynamic motion between bins
- Partially busy bit: 1 if exactly one buddy in use (0 if both/neither in use)
- Example using one group of eight pages
 - view also as two groups of four, four groups of two, or eight single pages
 - initially appears as a single chunk in 8-page free list
 - all other free lists are empty
 - all partially free bits are 0

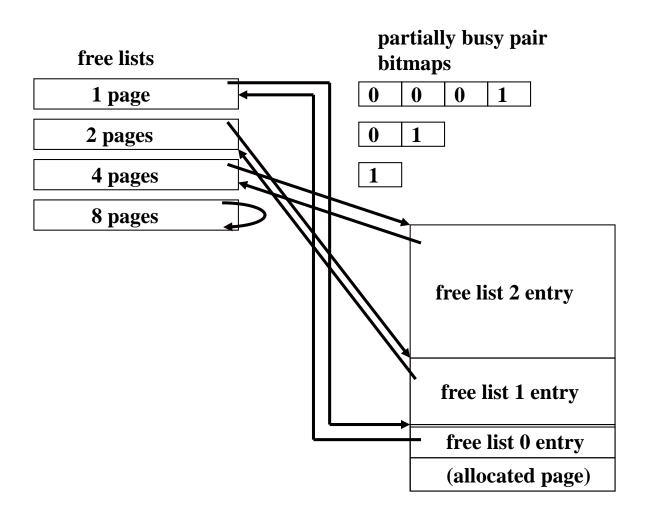
Initial configuration



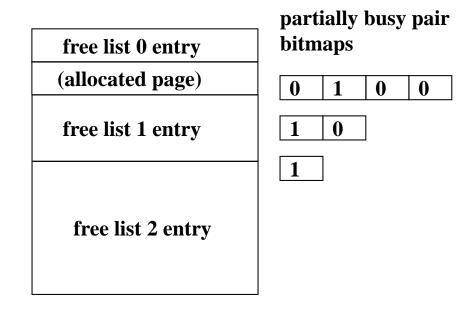
Allocation

- try the correct size free list
- if empty, try the next larger size and break up a chunk

- Request one page of order 0
 (order is log of # of pages, i.e., one page here)
 - any in free list 0 (1 page)? no…
 - any in free list 1 (2 pages)? no…
 - any in free list 2 (4 pages)? no…
 - any in free list 3 (8 pages)? yes! split it up recursively...



- Deallocation
 - check if buddy is free (is pair bit = 1?)
 - if both free, merge and check again (recursively)
- initial configuration for free example



- Free element (order is 0)
 - check (& flip) buddy bit in order 0
 - buddy was free → remove buddy from free list and merge (address remains the same for now)
 - check (& flip) buddy bit in order 1
 - buddy was free → remove buddy from free list and merge (address changes to start of first 4-page block)
 - repeat for order 2, then done; end with initial (all free) configuration