# ECE391 Computer System Engineering Lecture 25

Dr. Zbigniew Kalbarczyk
University of Illinois at Urbana- Champaign

Fall 2022

# Lecture Topics

Linux abstraction on device drivers

Next: Driver design process: I-mail design

#### **Aministrivia**

#### MP3.5 Code cutoff:

- -6:00 pm, Sunday, December 4
- Hand-in/demo (all group members should be present)
- Sunday (12/4), Monday (12/5), and Tuesday (12/6)

### Time slots posted at:

https://docs.google.com/spreadsheets/d/1tA7q5zk303KAQcpilaxGlfpNFfmJdn6AAWeKzVBs-zw/edit#gid=0

# Aministrivia

# Class Competition

- Thursday December 8
- More details to be provided

#### Final Exam:

- Tuesday December 13
- Time: 7:00 pm 10:00 pm

# Kernel Role in the System (1)

#### Process management

- implements the abstraction of several processes executing on top of a single/ multiple CPU(s)
- e.g., the kernel is responsible for creating/destroying and scheduling processes

#### Memory management

 supports virtual memory for all processes on top of the limited available resources

### Filesystem(s)

- supports different file systems
- almost everything in Linux can be treated as a file

# Kernel Role in the System (2)

#### Device control

- supports the system interactions with hardware devices
- almost every system operation eventually maps to a physical device
- device control operations are defined/implemented by a device driver specific to a given hardware device

#### Networking

- manages delivering data packets across program and network interfaces,
- handles routing and address resolution

#### **Device Drivers**

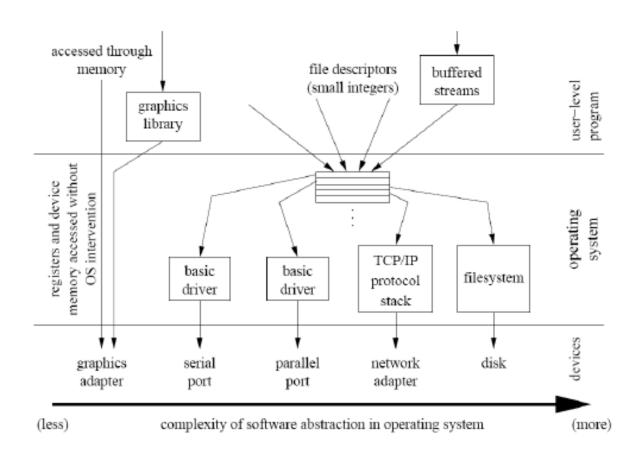
- Kernel interacts with I/O devices by means of device drivers
  - operate at the kernel level
  - include data structures and functions that control one or more devises, e.g., keyboard, monitors, network interfaces

#### Advantages:

- device code can be encapsulated in a module
- HW device responds to a well-defined programming interface
- hide the details on how the device works
- dynamic load/unload of device drivers

# Example of Range of Device Driver Complexity

Complexity grows from left to right



# Range of Device Driver Complexity

#### Leftmost example

- no driver used (other than for obtaining/release access rights)
- reads/writes go directly to device without OS intervention
- lack of system calls improves performance
- often used for graphics, and often with user-level libraries

#### Remaining examples use file descriptor abstraction

- sometimes with user-level libraries
- all reduced to using file descriptors to do system calls
- operating system
  - looks up file in array
  - uses file operations structure to hand off system call appropriately

# Range of Device Driver Complexity

- Simple drivers such as serial port/parallel port
- Drivers for network adapters add state for protocol, kernel-side buffering

- Drivers for disks support block transfer, read ahead, ....
  - usually used indirectly by file systems (inside kernel)
  - some user code (e.g., databases) uses raw disk commands

# Kernel Abstraction for Devices – Device Files

- I/O devices are treated as special files called device files
  - same system calls used to interact with files on disk can be used to work with I/O devices
  - e.g., same write() used to
    - write to a regular file or
    - send data to a printer by writing to /dev/lp0 device file
- According to the characteristics of the underlying device drivers, device files can be of two types
  - block
  - character

# Kernel Abstractions for Devices: Block Devices

- Block devices (underlies file systems)
  - data accessible only in blocks of fixed size, with size determined by device
  - can be addressed randomly
  - transfers to/from device are usually buffered (to) and cached (from) for performance
  - examples: disks, CD ROM, DVD

# Kernel Abstractions for Devices: Character Devices

#### Character device:

- Almost everything else, except network cards
  - Network cards are not directly associated with device files

#### A more constructive definition

- contiguous space (or spaces) of bytes
- some allow random access (e.g., magnetic tape driver)
- others available only sequentially (e.g., sound card)
- examples: keyboard, terminal, printers

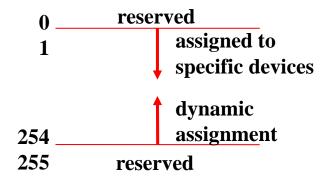
# Kernel Abstractions for Devices

- A device file is usually stored as a real file
  - Inode includes an identifier of the hardware device corresponding to the character or block device file

- Devices identified by major and minor numbers (traditionally both 8-bit)
  - major number is device type (e.g., specific model of disk, not just "disk")
  - minor number is instance number (if driver allows you to have more than one of a given model attached to your computer)

# Kernel Abstractions for Devices

- Major numbers each have
  - associated fops structure
  - name (see /proc/devices)



#### **Example of device files:**

Name	Type	Major	Minor	Description
/dev/hda	block	3	0	First IDE disk
/dev/hda2	block	3	2	Secondary primary partion of first IDE disk
/dev/tty0	char	3	0	Terminal

# Device Driver Registration

- Each system call issued on a device file is translated by the kernel into an invocation of a function of a corresponding device driver
- To achieve this, a device driver must register itself
  - A device driver statically compiled in the kernel is registered during the kernel initialization phase
  - A device driver compiled as a kernel module is registered when the module is loaded
  - insmod e.g., insmode ./hello.ko #link the module to the running kernel
  - rmmod e.g., rmmod ./hello #unlink the module from the running kernel

# Kernel Abstractions for Devices

Registering and unregistering a device (see linux/fs.h)

- to request a specific major #, pass it as input argument
  - returns 0 on success
  - returns negative value on failure
- for a dynamically assigned major #, pass 0 as input argument (major)
  - returns assigned major # on success (not 0!)
  - returns negative value on failure

# Kernel Abstractions for Devices

- both parameters must match those of registration call
- returns 0 on success
- returns negative value on failure

# Important Data Structures Used by the Driver

 Driver operations involve use of three kernel data structures:

#### file operations

jump table of functions in the driver that implement specific file operations

#### file structure (or file object)

- represents an open file
- no image in the permanent storage
- created by the kernel on "open" system call

#### inode

- used internally by the kernel to represent file metadata
- store in the permanent storage

# File Operations

#### File operations structure

- jump table of file operations / character driver operations
- generic instance for files on disk
- distinct instances for sockets, etc.
- one instance per device type

# Example: File Operation Structure in Imail

static struct file\_operations Imail\_fops = {

```
owner: THIS_MODULE,
read: Imail_read,
write: Imail_write,
poll: Imail_poll,
ioctl: Imail_ioctl,
release: Imail_release,
fsync: Imail_fsync,
};
```

# Example of module\_init() and module\_exit() for Imail

Kernel module must define two functions:

```
module_init (Imail_init) # invoked when the module is loaded into the kernel

module_exit (Imail_exit) # invoked when the module is removed from the kernel
```